# LINX(7) manual page

## Name

linx - LINX inter-process communication protocol

## Synopsis

**#include <linx.h>**
**#include <linx_socket.h>**
**#include <linx_ioctl.h>**
**#include <linx_types.h>**
**#include <sys/socket.h>**
**#include <sys/ioctl.h>**

## Description

**LINX** is a location transparent inter-process communication protocol. It is based on the message passing technology used in the Enea OSE family of real time operating systems.

LINX consists of a set of kernel modules and provides a standard socket based interface using its own protocol family, **PF_LINX.** There is also a LINX library that provides a more advanced messaging API. Applications normally access LINX through this library, but sometimes direct access via the socket interface may be necessary.

## Linx Concepts

An application that wants to communicate using LINX must first create a **LINX endpoint.** One thread may own multiple LINX endpoints simultaneously. A LINX endpoint is created with a non-unique name (a string). A handle which is used to refer to the endpoint in subsequent calls is returned to the application.

LINX endpoints communicate by sending and receiving **LINX signals.** The properties of LINX signals are described below.

Each LINX endpoint has a binary identifier called a **spid** which is unique within the node. A sending endpoint must know the spid of the destination endpoint to be able to communicate. The spid of a peer LINX endpoint is normally obtained by **hunting** for its name. When an endpoint with a matching name is found or created, LINX sends a hunt signal back to the hunting endpoint that appears to have been sent from the found endpoint. The spid can thus be obtained by looking at the sender of this signal.

A LINX endpoint can supervise, or **attach** to, the spid of a peer endpoint in order to be notified by a LINX signal when it is terminated or becomes unreachable.

The communication path between two LINX nodes is called a **LINX link.** A LINX link is created with a name string, and the same link may have different names on the opposite nodes, i.e. the link between nodes A and B may be called "LinkToB" on A, and "LinkToA" on B.

When hunting for an endpoint on a remote node, the name of the endpoint is prepended with the path of link names needed to reach the node, e.g. "LinkToB/LinkToC/EndpointName". LINX will create a virtual endpoint that acts as a local representation of the remote endpoint. LINX signals sent to the spid of a virtual endpoint are automatically routed to the proper destination on the remote node.

When a LINX endpoint is closed, its owned LINX signals are freed.

It is not allowed to use a LINX endpoint from two contexts at the same time. When a Linux process has multiple threads, it is not allowed to access a LINX endpoint from other contexts than the one that opened it. When **fork**(2) is called, the child process inherits copies of the parents socket related resources, including LINX endpoints. In this case, either the parent or the child shall close its LINX endpoints. A LINX endpoint is not removed until it has been closed by all of its owners.

# LINX Application Interfaces

LINX provides both a standard socket interface and a more advanced **LINX API** which is available through the LINX library, to be linked with the application. The LINX API is the recommended way for applications to access LINX, since it simplifies for the programmer by abstracting the direct socket interactions. It implements a set of functions specified in **linx.h**(3) .

The LINX socket interface is the underlying socket implementation provided by the LINX kernel module and is used by the LINX library. It is described in detail below.

It is possible for applications to use a combination of the LINX API and the LINX socket interface. In this case, the LINX API function **linx_get_descriptor**(3) can be used to obtain the socket descriptor of a LINX endpoint. This descriptor can be used together with other file descriptors in generic **poll**(2) or **select**(2) calls. Note that it is NOT allowed to call **close**(2) on a LINX socket descriptor obtained by a **linx_get_descriptor**(3) call.

# LINX Signals

A LINX signal consists of a mandatory leading 4-byte signal number, optionally followed by data. Thus, the size of a LINX signal buffer must be at least 4 bytes. LINX signal numbers are of type **LINX_SIGSELECT.** Signal numbers are mainly defined by the applications, but a few values are not allowed. Zero (0) is illegal and must not be used and 250-255 are reserved by LINX.

LINX provides endian conversion of the signal number if needed when a signal is sent to a remote node. The signal data is not converted.

# LINX Socket Interface

The LINX socket interface allows application programmers to access LINX using standard socket calls. It should be noted that only a subset of the socket calls are implemented and that additional features have been made available through ioctl calls. These deviations and features are described below.

**struct sockaddr_linx**
> When using the socket interface directly, a LINX endpoint is represented by a **sockaddr_linx** structure:

```
struct sockaddr_linx
{
        sa_family_t family;
        LINX_SPID   spid;
};
```

**family** shall be **AF_LINX** and **spid** is the spid of the LINX endpoint. The **sockaddr_linx** structure is type-casted into a generic **sockaddr** structure when passed to the socket interface function calls.

The following calls are provided by the LINX socket interface:

**socket()**

A LINX socket is created by calling the **socket**(2) function as:

**linx_sd = socket**(*PF_LINX*, *SOCK_DGRAM*, **0**)*;*

When a LINX socket is created, its name is unspecified. To assign a name to the socket, use the **LINX_IOCTL_HUNTNAME** ioctl request.

On success, the return value is a descriptor referencing the socket. On error, -1 is returned and *errno* is set to one of the following values:

**EPROTONOTSUPPORTED**
The protocol type is not supported. Only **PF_LINX** is accepted.

**ESOCKTNOSUPPORT**
The socket type is not supported. Only **SOCK_DGRAM** is accepted.

**ENOMEM**
Insufficient memory is available. Alternatively, the maximum number of spids has been reached. This value is configurable as a parameter to the LINX kernel module, see LINX Users Guide for more information.

**sendto()**

The **sendto**(2) function is called as:

**len = sendto**(*linx_sd*, *payload*, *size*, **0**, (**struct** *sockaddr***\***) **&***sockaddr_linx*, **sizeof**(**struct** *sockaddr_linx*))**;**

The *payload* shall be a LINX signal buffer and *size* shall be its length in bytes. Note that it is mandatory for a LINX signal to have a leading 4 byte signal number. The *spid* field of the *sockaddr_linx* structure shall be the spid of the destination endpoint.

On success, the number of bytes sent is returned. On error, -1 is returned and *errno* is set to one of the following values:

**EBADF**
An invalid descriptor was specified.

**ECONNRESET**
The destination endpoint has been killed.

**EINVAL**
Invalid argument passed.

**ENOMEM**
Insufficient memory is available.

**EOPNOTSUPP**
The sending LINX socket has not been assigned a name.

**EPIPE**
This error is reported at an attempt to send to the spid of a LINX endpoint that is being closed as the call occurs.

**sendmsg()**

The **sendmsg**(2) function is called as:

**len = sendmsg(***linx_sd***, ****msg***, 0);**

*msg* is a **msghdr** structure as defined in **sendmsg**(2) . The *msg_iov* field of the msghdr structure shall point to an **iovec** structure containing the LINX signal buffer to transmit and the *msg_iovlen* field shall be set to 1. Note that it is mandatory for a LINX signal to have a leading 4 byte signal number. The *msg_name* field shall be a pointer to a **sockaddr_linx** structure containing the spid of the destination endpoint and the *msg_namelen* field shall be set to the size of the sockaddr_linx structure. The ancillary fields and the flags field of the msghdr structure shall not be used and be set to zero.

On success, the number of bytes sent is returned. On error, -1 is returned and *errno* is set to one of the following values:

**EBADF**
An invalid descriptor was specified.

**ECONNRESET**
The destination endpoint has been killed. Note that this case is accepted and the signal is silently discarded by LINX.

**EINVAL**
Invalid argument passed.

**ENOMEM**
Insufficient memory is available.

**EOPNOTSUPP**
The sending socket has not been assigned a name.

**EPIPE**
This error is reported at an attempt to send to the spid of a LINX endpoint that is being closed as the call occurs.

**recvfrom()**

The **recvfrom**(2) function is called as:

**len = recvfrom(***linx_sd***, *payload*, *size*, 0, (struct** *sockaddr***)** **&***sockaddr_linx***, sizeof(struct** *sockaddr_linx***));**

It is used to receive any LINX signal from any LINX endpoint. It can not be used when signal number filtering and/or sender filtering is needed, see **recvmsg**(2) . The first signal in the sockets receive queue is returned in the supplied *payload* buffer. If no signal is currently available at the socket, the call blocks until a signal is received. The sender of the signal is returned in the **sockaddr_linx** structure. Note that the *size* of the payload buffer must be at least 4 bytes, since it is mandatory for a LINX signal to have a 4 byte leading signal number.

On success, the number of bytes received is returned. If the received signal is larger than the supplied payload buffer, zero is returned and the signal buffer size is written as a 32-bit value in the first 4 bytes of the payload buffer. On error, -1 is returned and *errno* is set to one of the following values:

**EBADF**
An invalid descriptor was specified.

**EFAULT**
Invalid payload buffer pointer provided.

**EINVAL**
Invalid argument passed.

**ENOMEM**
Insufficient memory is available.

**EOPNOTSUPP**
The receiving socket has not been assigned a name.

**recvmsg()**
The **recvmsg**(2) function is called as:

**len = recvmsg(***sd***, ***msg***, 0);**

*msg* is a pointer to a **msghdr** structure as defined in **recvmsg**(2) . A LINX signal buffer shall be supplied in an **iovec** structure pointed to from the *msg_iov* field and the **msg_iovlen** field shall be set to 1. Note that the size of the supplied buffer must be at least 4 bytes, since it is mandatory for a LINX signal to have a 4 byte leading signal number.

The **recvmsg**(2) call supports signal number filtering and sender filtering. This allows the user to specify which signal numbers shall be received and/or from which sender. The signal filter is described by a **linx_receive_filter_param** structure:

```
struct linx_receive_filter_param
{
        LINX_SPID            from;
        LINX_OSBUFSIZE       sigselect_size;
        const LINX_SIGSELECT *sigselect;
};
```

The **from** field specifies that only signals from a specific spid should be received and **sigselect** is an array of **LINX_SIGSELECT** numbers to be received. The first position in the array contains the number of entries in the list that follows. If the first position is set to a negative count, all LINX signals except those listed will be received. The size of the array is **sigselect_size.**

The filtering uses ancillary fields in the **msghdr** structure and is described by the following example code:

```
struct msghdr msg;
char cmsg[CMSG_SPACE(sizeof(struct linx_receive_filter_param))];
struct linx_receive_filter_param * rfp;
struct iovec iov;
const LINX_SIGSELECT sig_sel[] = { 1, EXAMPLE_SIG };
```

```
rfp = ((struct linx_receive_filter_param *)
        (CMSG_DATA(((struct cmsghdr *)cmsg)))));
rfp->sigselect_size = sizeof(sig_sel);
rfp->from = from_spid;
rfp->sigselect = sig_sel;
msg.msg_name = (void*)&linx_addr;
msg.msg_namelen = sizeof(struct sockaddr_linx);
msg.msg_iov = &iov;
msg.msg_iovlen = 1;
msg.msg_flags = 0;
msg.msg_control = cmsg;
msg.msg_controllen =
        CMSG_SPACE(sizeof(struct linx_receive_filter_param));
((struct cmsghdr *)cmsg)->cmsg_len = msg.msg_controllen;
((struct cmsghdr *)cmsg)->cmsg_level = 0;
((struct cmsghdr *)cmsg)->cmsg_type = 0;
iov.iov_base = *signal;
iov.iov_len = sigsize;
read_size = recvmsg(linx->socket, &msg, 0);
```

On success, the number of bytes received is returned. If the received signal is larger than the supplied payload buffer, zero is returned and the signal buffer size is written as a 32-bit value in the first 4 bytes of the payload buffer. On error, -1 is returned and *errno* is set to one of the following values:

**EBADF**
An invalid descriptor was specified.

**EFAULT**
An invalid msghdr structure was provided.

**EINVAL**
Invalid argument passed.

**ENOMEM**
Insufficient memory is available.

**EOPNOTSUPP**
The sending socket has not been assigned a name.

**poll()**

LINX socket descriptors can be used in the **poll**(2) call.

The call returns a bitmask that indicates the state of the LINX socket receive queues. The following possible bits can be set at return from this function: **POLLERR** if the LINX socket is in an error state, **POLLHUP** if the LINX socket has been shutdown/released, **POLLIN** and **POLLRDNORM** if the LINX socket has data to read in receive queue.

**select()**

LINX socket descriptors can be used together with other file descriptors in the **select**(2) call.

**ioctl()**

IOCTL requests are sent to a LINX socket using the **ioctl**(2) call. See below for all IOCTL request codes supported by LINX sockets.

**close()**

A LINX socket created by a *socket*(2) call can be closed with *close*(2) . Note that this function shall NOT be used on any LINX socket descriptor created with *linx_open*(3) or obtained by the LINX API function *linx_get_descriptor*(3) .

On success, 0 is returned, otherwise -1 is returned and *errno* can be one of the following errors:

**EBADF**
An invalid descriptor was specified.

Only the calls described above are supported by a LINX socket. The following are NOT supported on LINX sockets and shall not be used: *bind*(2) , *connect*(2) , *socketpair*(2) , *accept*(2) , *getname*(2) , *listen*(2) , *shutdown*(2) , *setsockopt*(2) , *getsockopt*(2) , *mmap*(2) and *sendpage*(2) .

# IOCTL Request Codes

The following IOCTL request codes can be accessed using *ioctl*(2) on LINX sockets:

**LINX_IOCTL_SEND**

Sends a signal from a LINX socket **sd**. The correct syntax is:

**struct linx_sndrcv_param \****sndrcv*;

*error* **= ioctl(***sd*, *LINX_IOCTL_SEND*, *sndrcv***);**

*sndrcv* is a **linx_sndrcv_param** structure:

```
struct linx_sndrcv_param
{
    __u32 from;
    __u32 to;
    __u32 size;
    __u32 sig_attr;
    __u32 sigselect_size;
    __u64 sigselect;
    __u32 tmo;
    __u64 buffer;
};
```

**from** is the spid of the sender, the sender does not need to be the current socket, the signal is sent from the current socket but the receiver will see the **from** as the sending spid, the default case is that **from** is the spid of the current LINX socket. **to** is the spid of the receiver, **size** is the size of the signal to be sent and **sig_attr** are the attributes of the signal. The **sigselect_size** and **sigselect** are not used by **LINX_IOCTL_SEND** and **notused** is for padding the struct since it needs to be 64-bit aligned but should be zeroed for future use. **buffer** is the pointer to the buffer to be sent, it is passed as a 64 bit unsigned value to be both 32-bit and 64-bit compatible.

On success, the number of bytes sent is returned.

**LINX_IOCTL_RECEIVE**

Receives a signal on a LINX socket **sd**. The call will block until a signal is received. The correct syntax is:

**struct linx_sndrcv_param \****sndrcv*;

```
error = ioctl(sd, LINX_IOCTL_RECEIVE, sndrcv);
```

*sndrcv* is a **linx_sndrcv_param** structure:

```
struct linx_sndrcv_param
{
    __u32 from;
    __u32 to;
    __u32 size;
    __u32 sig_attr;
    __u32 sigselect_size;
    __u64 sigselect;
    __u32 tmo;
    __u64 buffer;
};
```

The **from** field should be set to **LINX_ILLEGAL_SPID** if signals from anyone should be received or if only signals from a specific spid should be received then **from** should be set to that spid. If **LINX_ILLEGAL_SPID** was set the **from** will contain the spid of the sender after the return of the call. The **to** field is not used when receiving a signal. The **size** field is the size of the provided buffer. When the call returns the **sig_attr** field is set to the attribute the signal carriers. The **sigselect** field is an array of **LINX_SIGSELECT** numbers to be received. The first position in the array contains the number of entries in the list that follows. If the first position is set to a negative count, all LINX signals except those listed will be received. The size of the array is **sigselect_size.** When the **tmo** field is used the call waits maximum **tmo** milliseconds before returning even if no signal has been received, if a blocking receive is requested the **tmo** field should be ~0 (0xFFFFFFFF), this will block forever. If no signal is received the. **buffer** pointer will be set to NULL (the provided buffer is always consumed). The **buffer** field is a pointer to the buffer provided by the user.

On success, the number of bytes received is returned. If the received signal is larger than the supplied payload buffer, zero is returned and the signal buffer size is written as a 32-bit value in the first 4 bytes of the payload buffer.

**LINX_IOCTL_REQUEST_TMO**

Request a timeout, a signal is sent to the requesting LINX endpoint when a timeout has expired. The correct syntax is:

**struct linx_tmo_param \****tmo_param***;**

*error* **= ioctl(***sd*, *LINX_IOCTL_REQUEST_TMO*, *tmo_param***);**

*tmo_param* is a **linx_tmo_param** structure:

```
struct linx_tmo_param
{
    LINX_OSTIME tmo;
    LINX_OSBUFSIZE sigsize;
    union LINX_SIGNAL *sig;
    LINX_OSTMOREF tmoref;
};
```

**tmo** is the timeout in milliseconds and the actual timeout time is rounded upward to the nest larger tick, the call guarantees at least the number of milliseconds requested, **sig** is a pointer to the signal that will be returned when the timeout expires, if **sig** is LINX_NIL then the default timeout signal with signal number LINX_OS_TMO_SIG is received instead, **sigsize** is the size of the provided signal, if no signal is provided the value must be set to zero.

On success, an timemout reference is returned in **tmoref**. This reference can be used in **LINX_IOCTL_CANCEL_TMO** and **LINX_IOCTL_MODIFY_TMO**.

**LINX_IOCTL_CANCEL_TMO**

Cancel a pending timeout, the correct syntax is:

**struct linx_tmo_param \****tmo_param**;*

*error* **= ioctl(***sd*, *LINX_IOCTL_CANCEL_TMO*, *tmo_param***);**

*tmo_param* is a **linx_tmo_param** structure:

```
struct linx_tmo_param
{
    LINX_OSTIME tmo;
    LINX_OSBUFSIZE sigsize;
    union LINX_SIGNAL *sig;
    LINX_OSTMOREF tmoref;
};
```

**tmo**, **sigsize** and **sig** are ignored, **tmoref** is used to identify which timeout is to be canceled, it is guaranteed that the timeout signal cannot be received after cancellation.

**LINX_IOCTL_MODIFY_TMO**

Modifies a pending timeout, the correct syntax is:

**struct linx_tmo_param \****tmo_param**;*

*error* **= ioctl(***sd*, *LINX_IOCTL_MODIFY_TMO*, *tmo_param***);**

*tmo_param* is a **linx_tmo_param** structure:

```
struct linx_tmo_param
{
    LINX_OSTIME tmo;
    LINX_OSBUFSIZE sigsize;
    union LINX_SIGNAL *sig;
    LINX_OSTMOREF tmoref;
};
```

**sigsize** and **sig** are ignored, **tmoref** is used to identify which timeout is to be modified, **tmo** is the new timeout value.

**LINX_IOCTL_REQUEST_NEW_LINK**

Request a signal when a new link is available, the correct syntax is:

**struct linx_new_link_param \****new_link_param*;

*error* **= ioctl(***sd*, *LINX_IOCTL_REQUEST_NEW_LINK*, *new_link_param***);**

*new_link_param* is a **linx_new_link_param** structure:

```
struct linx_new_link_param
{
    uint32_t token;
    uint32_t new_link_ref;
};
```

**token** is passed to and from the LINX kernel module keeping track of which links the caller already have been notified about, the **token** value is ignored the first time a LINX endpint requests a new link signal. The **token** received in the new link signal should then be used in the next new link signal request. **new_link_ref** is used to cancel a pending new link signal request.

The syntax of the new link signal received when a new link is available is:

```
struct linx_new_link {
    LINX_SIGSELECT signo;
    LINX_NLTOKEN token;
    int name;
    int attr;
    char buf[1];
};
```

**signo** is LINX_OS_NEW_LINK_SIG, **token** is the token value to be used in the next request for a new link signal, **name** is the offset into **buf** where the name of the new link is stored, the name is null terminated, **attr** is the offset into **buf** where the attributes, if any, of the link are stored, the attribute string is null terminted, **buf** is a character buffer containg the name and the attributes, if any, of the link.

**LINX_IOCTL_CANCEL_NEW_LINK**

Cancels a pending new link signal request, the correct syntax is:

**struct linx_new_link_param \****new_link_param*;

*error* **= ioctl(***sd*, *LINX_IOCTL_CANCEL_NEW_LINK*, *new_link_param***);**

*new_link_param* is a **linx_new_link_param** structure:

```
struct linx_new_link_param
{
    uint32_t token;
    uint32_t new_link_ref;
};
```

**token** is ignored, **new_link_ref** is used to identify which pending new link signal request is to be cancelled, after the cancellation it is guarnateed that no more new link signals can be received.

## LINX_IOCTL_HUNTNAME

Sets the name of a LINX socket **sd** and returns its binary LINX endpoint identifier, the spid. The correct syntax is:

**struct linx_huntname *** *huntname*;

*error* **= ioctl(** *sd*, *LINX_IOCTL_HUNTNAME*, *huntname* **);**

*huntname* is a **linx_huntname** structure:

```
struct linx_huntname
{
        LINX_SPID spid;
        size_t    namelen;
        char     *name;
};
```

**namelen** is the size in bytes of the string **name** that contains the name to assign to the socket. On success, **spid** is set to the spid assigned to the LINX socket.

## LINX_IOCTL_HUNT

Hunts for a LINX endpoint (that has been assigned a name) to obtain its spid. The correct syntax is:

**struct linx_hunt_param *** *hunt_param*;

*error* **= ioctl(** *sd*, *LINX_IOCTL_HUNT*, *hunt_param* **);**

*hunt_param* is a **linx_hunt_param** structure:

```
struct linx_hunt_param
{
        LINX_OSBUFSIZE     sigsize;
        union LINX_SIGNAL *sig;
        LINX_SPID          from;
        size_t             namelen;
        char              *name;
```

```
                };
```

The **sig** parameter optionally holds a signal of size **sigsize** to be received when the other LINX socket is available. If no signal (NULL) is provided, the LINX default hunt signal of type **LINX_OS_HUNT_SIG** will be used. The **from** parameter shall be set to the owner of the hunt. In the normal case, this is the spid of the LINX socket performing the hunt call. If the spid associated with a different LINX socket is provided, the hunt can be cancelled by closing that socket. The hunt signal is always sent to the LINX socket performing the hunt call. The **namelen** is the size of the string **name** that contains the name of the LINX endpoint to be hunted for.

## LINX_IOCTL_ATTACH

Attaches to a LINX endpoint in order to supervise it, i.e. to get an attach signal if it becomes unavailable. The correct syntax is:

```
struct linx_attach_param *attach_param;

error = ioctl(sd, LINX_IOCTL_ATTACH, attach_param);
```

*attach_param* is a **linx_attach_param** structure:

```
        struct linx_attach_param
        {
                LINX_SPID          spid;
                LINX_OSBUFSIZE     sigsize;
                union LINX_SIGNAL *sig;
                LINX_OSATTREF      attref;
        };
```

The **spid** field is the spid of the LINX endpoint to supervise. The **sig** parameter optionally holds a LINX signal of size **sigsize** to be received when the supervised LINX endpoint becomes unavailable. If no signal (NULL) is provided, the LINX default attach signal of type **LINX_OS_ATTACH_SIG** will be used.

On success, an attach reference is returned in **attref.** This reference can be used in **LINX_IOCTL_DETACH** later.

## LINX_IOCTL_DETACH

Detaches from a supervised LINX endpoint, i.e. stops supervising it. The correct syntax is:

```
struct linx_detach_param *detach_param;

error = ioctl(sd, LINX_IOCTL_DETACH, detach_param);
```

The *detach_param* parameter is a **struct linx_detach_param** with the following fields:

```
struct linx_detach_param
{
        LINX_OSATTREF  attref;
}
```

The **attref** field is an attach reference returned from a **LINX_IOCTL_ATTACH** call.

## LINX_IOCTL_SET_RECEIVE_FILTER

Sets up a receive filter prior to a **select**(2) call. The correct syntax is:

**struct linx_receive_filter_param \****rfp*;

*error* **= ioctl(***sd***,** *LINX_IOCTL_SET_RECEIVE_FILTER***,** *rfp***);**

*rfp* is a **linx_receive_filter_param** structure:

```
struct linx_receive_filter_param
{
        LINX_SPID               from;
        LINX_OSBUFSIZE          sigselect_size;
        const LINX_SIGSELECT *sigselect;
};
```

The **from** parameter specifies that only signals from a specific spid should be received and **sigselect** is an array of **LINX_SIGSELECT** numbers to be received. The first position in the array contains the number of entries in the list that follows. If the first position is set to a negative count, all LINX signals except those listed will be received. The size of the array is **sigselect_size.**

## LINX_IOCTL_REGISTER_LINK_SUPERVISOR

This command is now obsolete; use the LINX_IOCTL_REQUEST_NEW_LINK instead.

## LINX_IOCTL_UNREGISTER_LINK_SUPERVISOR

This command is now obsolete; use the LINX_IOCTL_CANCEL_NEW_LINK instead.

## LINX_IOCTL_VERSION

Returns the version of the LINX kernel module. The correct syntax is:

**unsigned int** *version*;

```
error = ioctl(sd, LINX_IOCTL_VERSION, &version);
```

On success, the **version** parameter contains the version of the LINX kernel module. The LINX version number is a 32-bit number composed of an 8-bit major version, an 16-bit minor version, and a 8-bit seq (patch) number.

## LINX_IOCTL_INFO

Retrieves information from the LINX kernel module. The correct syntax is:

```
struct linx_info info;

error = ioctl(sd, LINX_IOCTL_INFO, &info);
```

The **info** parameter is a **struct linx_info** with the following fields:

```
struct linx_info
{
        int    type;
        void  *type_spec;
};
```

The **type** field indicates the requested type of information and **type_spec** is a pointer to a struct that will contain input and return parameters.

Note that information retrieved with LINX_IOCTL_INFO may have become inaccurate when used in a subsequent call. The application must be prepared to handle errors related to this.

The different kinds of information that can be retrieved from the LINX kernel module are:

### LINX_INFO_SUMMARY

Provides a summary of the most important information from the LINX kernel module.

```
struct linx_info info;

struct linx_info_summary info_summary;

info.type = LINX_INFO_SUMMARY;

info.type_spec = &info_summary;
```

The **linx_info_summary** structure is defined as:

```
struct linx_info_summary
{
        int no_of_local_sockets;
        int no_of_remote_sockets;
        int no_of_link_sockets;
        int no_of_pend_attach;
        int no_of_pend_hunt;
        int no_of_queued_signals;
};
```

The **no_of_local_sockets** field is the number of LINX sockets open locally, **no_of_remote_sockets** is the number of internal sockets open that have been created by the LINX kernel module to represent remote LINX endpoints. The **no_of_link_sockets** field is the number of open sockets representing LINX links to other nodes. The **no_of_pend_attach** field is the number of pending attaches, **no_of_pend_hunt** is the number of pending hunts and **no_of_queued_signals** is the number of queued signals.

**LINX_INFO_SOCKETS**

Returns the number of open LINX sockets and their LINX endpoint identifiers (spids).

**struct linx_info** *info;*

**struct linx_info_sockets** *info_sockets;*

**info.type =** *LINX_INFO_SOCKETS;*

**info.type_spec =** *&info_sockets;*

The **linx_info_sockets** structure is defined as:

```
struct linx_info_sockets
{
        LINX_OSBOOLEAN local;
        LINX_OSBOOLEAN remote;
        LINX_OSBOOLEAN link;
        int            buffer_size;
        int            no_of_sockets;
        LINX_SPID      *buffer;
};
```

If **local** is true, local sockets are included in the output, if **remote** is true, remote sockets are included and if **link** is true, sockets representing LINX links are included. The number of LINX sockets matching the search is returned in **no_of_sockets** and the array of spids is returned in the provided **buffer**. of size **buffer_size** bytes. If the provided buffer is too small, not all sockets will be included.

### LINX_INFO_TYPE

Returns the type of a LINX endpoint.

**struct linx_info** *info*;

**struct linx_info_type** *info_type*;

**info.type =** *LINX_INFO_TYPE*;

**info.type_spec =** *&info_type*;

The **linx_info_type** structure is defined as:

```
struct linx_info_type
{
        LINX_SPID  spid;
        int        type;
};
```

The **spid** field is the identifier of the LINX endpoint for which the type is requested. The type is returned in **type**. A LINX endpoint can be of types: **LINX_TYPE_UNKNOWN**, **LINX_TYPE_LOCAL**, **LINX_TYPE_REMOTE**, **LINX_TYPE_LINK**, **LINX_TYPE_ILLEGAL** or **LINX_TYPE_ZOMBIE**.

### LINX_INFO_STATE

Returns the state of a LINX endpoint.

**struct linx_info** *info*;

**struct linx_info_state** *info_state*;

**info.type =** *LINX_INFO_STATE*;

```
info.type_spec = &info_state;
```

The **linx_info_state** structure is defined as:

```
struct linx_info_state
{
        LINX_SPID  spid;
        int        state;
};
```

The **spid** field is the identifier of the LINX endpoint for which the state is requested. The state is returned in **state**. A LINX socket can be in states: **LINX_STATE_UNKNOWN**, **LINX_STATE_RUNNING**, **LINX_STATE_RECV** or **LINX_STATE_POLL**.

**LINX_INFO_FILTERS**

Returns information about the receive filters set up by a LINX endpoint.

**struct linx_info** *info*;

**struct linx_info_filters** *info_filters*;

**info.type =** *LINX_INFO_FILTERS*;

**info.type_spec =** *&info_filters*;

The **linx_info_filters** structure is defined as:

```
struct linx_info_filters
{
        LINX_SPID       spid;
        LINX_SPID       from_filter;
        int             buffer_size;
        int             no_of_sigselect;
        LINX_SIGSELECT *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint for which the receive filter is requested. If the endpoint has setup a receive filter only accepting

signals from a specific LINX endpoint, the spid of that endpoint is returned in **from_filter**. The number of LINX_SIGSELECT signal numbers in the receive filter is returned in **no_of_sigselect**. A copy of the filter of the LINX endpoint is returned in **buffer**. The first element of the filter is the amount of LINX_SIGSELECT that follow it.

**buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all signal numbers in the filter are included. It is possible to know whether all LINX_SIGSELECT were copied to the **buffer** array by comparing its first element with **no_of_sigselect**.

### LINX_INFO_RECV_QUEUE

Returns the receive queue of a LINX endpoint.

```
struct linx_info info;

struct linx_info_recv_queue info_recv_queue;

info.type = LINX_INFO_RECV_QUEUE;

info.type_spec = &info_recv_queue;
```

The **linx_info_recv_queue** structure is defined as:

```
struct linx_info_recv_queue
{
        LINX_SPID               spid;
        int                     buffer_size;
        int                     no_of_signals;
        struct linx_info_signal *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint for which the receive queue is requested. The number of signals in the queue is returned in **no_of_signals**. An array with queue information is returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all signals are included. The **linx_info_signal** structure is defined as:

```
struct linx_info_signal
{
        LINX_SIGSELECT signo;
        int            size;
        LINX_SPID      from;
};
```

The **signo** field holds the signal number, **size** is the size in bytes of the signal and **from** is the spid of the sending LINX endpoint.

### LINX_INFO_RECV_QUEUE_2

Returns the receive queue of a LINX endpoint, with information about OOB messages in queue.

**struct linx_info** *info;*

**struct linx_info_recv_queue_2** *info_recv_queue;*

**info.type =** *LINX_INFO_RECV_QUEUE_2;*

**info.type_spec =** *&info_recv_queue;*

The **linx_info_recv_queue** structure is defined as:

```
struct linx_info_recv_queue_2
{
        LINX_SPID  spid;
        int        buffer_size;
        int        no_of_signals;
        char       *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint for which the receive queue is requested. The number of signals in the queue is returned in **no_of_signals**. An array with queue information is returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all signals are included. The **linx_info_signal_2** structure is defined as:

```
struct linx_info_signal_2
{
        LINX_SIGSELECT signo;
        int            size;
        LINX_SPID      from;
        uint32_t       flags;
};
```

The **signo** field holds the signal number, **size** is the size in bytes of the signal and **from** is the spid of the sending LINX endpoint. **flags** specifies the type of transmission. Values of this argument are formed by logically OR'ing zero or more of the following flags:

**MSG_OOB** A signal sent out-of-band.

### LINX_INFO_PEND_ATTACH

Returns information about pending attaches **from** or **to** a LINX endpoint.

**struct linx_info** *info;*

**struct linx_info_pend_attach** *info_pend_attach;*

**info.type =** *LINX_INFO_PEND_ATTACH;*

**info.type_spec =** *&info_pend_attach;*

The **linx_info_pend_attach** structure is defined as:

```
struct linx_info_pend_attach
{
        LINX_SPID               spid;
        int                     from_or_to;
        int                     buffer_size;
        int                     no_of_attaches;
        struct linx_info_attach *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint for which attach information is requested. If **from_or_to** is set to **LINX_ATTACH_FROM,** information about attaches from the spid is returned. If it is set to **LINX_ATTACH_TO,** information about attaches to the spid is returned. The number of attaches to/from the **spid** is returned in **no_of_attaches**, Information about the attaches are returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all attaches are included. The **linx_info_attach** structure is defined as:

```
struct linx_info_attach
{
        LINX_SPID               spid;
        LINX_OSATTREF           attref;
        struct linx_info_signal attach_signal;
};
```

The **spid** is the identifier of the LINX endpoint that has attached or has been attached to (depending on what **from_or_to** is set to). The **attref** field is the attach reference and **attach_signal** is the attach signal.

## LINX_INFO_PEND_HUNT

Returns information about pending hunts issued from any LINX endpoint.

```
struct linx_info info;

struct linx_info_pend_hunt info_pend_hunt;

info.type = LINX_INFO_PEND_HUNT;

info.type_spec = &info_pend_hunt;
```

The **linx_info_pend_hunt** structure is defined as:

```
struct linx_info_pend_hunt
{
        LINX_SPID               spid;
        int                     buffer_size;
        int                     strings_offset;
        int                     no_of_hunts;
        struct linx_info_hunt *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint for which hunt information is requested. The number of pending hunts is returned in **no_of_hunts** and information about each pending hunt is returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all hunts are included. The **strings_offset** is the offset into the **buffer** where the name strings are stored. Each **linx_info_hunt** structure is defined as:

```
struct linx_info_hunt
{
        struct linx_info_signal hunt_signal;
        LINX_SPID               owner;
        char                    *hunt_name;
};
```

The **owner** field is the owner of the pending hunt and **hunt_name** is a string containing the name hunted for. The **hunt_signal** is the hunt signal. The **linx_info_signal** structure is described under **LINX_INFO_RECV_QUEUE**.


### LINX_INFO_PEND_TMO


Returns information about pending timeouts issued from any LINX endpoint.


```
struct linx_info info;

struct linx_info_pend_tmo info_pend_tmo;

info.type = LINX_INFO_PEND_TMO;

info.type_spec = &info_pend_tmo;
```


The **linx_info_pend_tmo** structure is defined as:


```
struct linx_info_pend_tmo
{
        LINX_SPID              spid;
        int                    buffer_size;
        int                    no_of_timeouts;
        struct linx_info_tmo  *buffer;
};
```


The **spid** field is the identifier of the LINX endpoint for which timeout information is requested. The number of pending timeouts is returned in **no_of_timeouts** and information about each pending timeout is returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the buffer is too small, not all timeouts are included. Each **linx_info_tmo** structure is defined as:


```
struct linx_info_tmo
{
        LINX_OSTIME            tmo;
        LINX_OSTMOREF          tmoref;
        struct linx_info_signal tmo_signal;
};
```


The **tmo** field is the remaining time and **tmoref** is the timeout reference. The **tmo_signal** holds the timeout signal information.

**LINX_INFO_SIGNAL_PAYLOAD**

Returns the payload of a signal owned by a LINX endpoint.

```
struct linx_info info;

struct linx_info_signal_payload info_signal_payload;

info.type = LINX_INFO_SIGNAL_PAYLOAD;

info.type_spec = &info_signal_payload;
```

The **linx_info_signal_payload** structure is defined as:

```
struct linx_info_signal_payload
{
        LINX_SPID               spid;
        int                     buffer_size;
        struct linx_info_signal signal;
        int                     payload_size;
        char                    *buffer;
};
```

The **spid** field is the identifier of the LINX endpoint owning the signal and **signal** is a **linx_info_signal** structure returned from a previous LINX_INFO call. The signal buffer will be returned in the provided **buffer**. The **buffer_size** is the size in bytes of the buffer. If the provided buffer is too small, only the beginning of the signal buffer is returned. The **payload_size** shows the size in bytes of the returned signal payload. If the provided buffer is larger than the signal payload, **payload_size** will be less than **buffer_size**. If no signal payload matching the **signal** the **payload_size** will be set to zero.

**LINX_INFO_NAME**

Returns the name of a LINX endpoint.

```
struct linx_info info;

struct linx_info_name info_name;

info.type = LINX_INFO_NAME;

info.type_spec = &info_name;
```

The **linx_info_name** structure is defined as:

```
struct linx_info_name
{
        LINX_SPID spid;
        int       namelen;
        char     *name;
};
```

The **spid** field is the identifier of the LINX endpoint for which the name is requested. The **namelen** field is the length of the provided **name** buffer in which the name is returned. If the LINX socket endpoint has not been assigned a name yet zero is returned and **name** is set to the empty string.

### LINX_INFO_OWNER

Returns the process (PID) that owns a LINX endpoint.

**struct linx_info** *info*;

**struct linx_info_owner** *info_owner*;

**info.type =** *LINX_INFO_OWNER*;

**info.type_spec =** *&info_owner*;

The **linx_info_owner** structure is defined as:

```
struct linx_info_owner
{
        LINX_SPID spid;
        pid_t    owner;
};
```

The **spid** field is the identifier of the LINX endpoint for which the information is requested. On success, the PID of the owning process is returned in **owner**.

### LINX_INFO_STAT

Returns statistics for a LINX endpoint. This requires that the LINX kernel module has been compiled with the "SOCK_STAT=yes" setting.

```
struct linx_info info;

struct linx_info_stat info_stat;

info.type = LINX_INFO_STAT;

info.type_spec = &info_stat;
```

The **linx_info_stat** structure is defined as:

```
struct linx_info_stat
{
        LINX_SPID spid;
        uint64_t no_sent_local_signals;
        uint64_t no_recv_local_signals;
        uint64_t no_sent_local_bytes;
        uint64_t no_recv_local_bytes;
        uint64_t no_sent_remote_signals;
        uint64_t no_recv_remote_signals;
        uint64_t no_sent_remote_bytes;
        uint64_t no_recv_remote_bytes;
        uint64_t no_sent_signals;
        uint64_t no_recv_signals;
        uint64_t no_sent_bytes;
        uint64_t no_recv_bytes;
        uint64_t no_queued_bytes;
        uint64_t no_queued_signals;
};
```

The **spid** is the identifier of the LINX endpoint for which statistics is required. If the LINX kernel module has not been compiled with "SOCK_STAT=yes", ioctl returns -1 and errno is set to ENOSYS.

# Known Bugs

None.

# See Also

Generic LINX for Linux man-page:
**linx**(7) (this document)

LINX API man-pages:
**linx.h**(3) , **linx_types.h**(3) ,

**linx_alloc**(3) , **linx_attach**(3) , **linx_cancel_tmo**(3) , **linx_close**(3) ,
**linx_detach**(3) , **linx_free_buf**(3) , **linx_free_name**(3) , **linx_free_stat**(3) ,
**linx_get_descriptor**(3) , **linx_get_name**(3) , **linx_get_spid**(3) ,
**linx_get_stat**(3) , **linx_hunt**(3) , **linx_hunt_from**(3) , **linx_modify_tmo**(3) ,
**linx_open**(3) , **linx_receive**(3) , **linx_receive_from**(3) ,
**linx_receive_w_tmo**(3) , **linx_request_tmo**(3) , **linx_send**(3) , **linx_send_w_opt**(3) ,
**linx_send_w_s**(3) , **linx_sender**(3) , **linx_set_sigsize**(3) , **linx_sigattr**(3) , **linx_sigsize**(3)

Related LINX applications
**linxcfg**(1) , **linxdisc**(8) , **linxdisc.conf**(5) , **linxstat**(1) , **mkethcon**(1) , **mklink**(1) , **mktcpcon**(1) ,
**rmethcon**(1) , **rmlink**(1) , **rmtcpcon**(1)

Related generic Linux man-pages:
**socket**(2) , **close**(2) , **sendto**(2) , **sendmsg**(2) , **recvfrom**(2) , **recvmsg**(2) , **poll**(2) , **select**(2) , **ioctl**(2)

# Author

Enea LINX team

# LINX_ALLOC(3) manual page

## Name

linx_alloc() - Allocate a signal buffer on a LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**union LINX_SIGNAL *linx_alloc(LINX *** *linx* **, LINX_OSBUFSIZE** *size* **, LINX_SIGSELECT** *sig_no* **);**

## Description

**linx_alloc**() is used to allocate a signal buffer of the indicated size The first 4 bytes of the signal buffer are initialized with the provided signal number *sig_no*. The signal number can be replaced later by writing a new value in the first 4 bytes of the buffer.

The new buffer is owned by the LINX endpoint with which it was created. The ownership is transferred to the receiving LINX endpoint when the signal is successfully sent using **linx_send**(3) . If the signal is not sent, the owner must free the buffer by calling **linx_free_buf**(3) using the same **linx** handle. The buffer will be freed if the owning LINX endpoint is closed or the process owning the endpoint exits.

*linx* is the handle to the LINX endpoint.

*size* is the size in bytes of the new signal buffer. The minimum size is 4 bytes, needed to store the signal number.

*sig_no* is the signal number stored at the beginning of the signal buffer.

## Return Value

On success, a pointer to the allocated signal buffer is returned. On error, LINX_NIL is returned and *errno* will be set.

## Errors

**EMSGSIZE** The *size* is invalid.

**ENOMEM** Insufficient memory is available.

## Bugs/Limitations

None.

## Notes

A LINX signal can be only be sent by the LINX endpoint that owns it, i.e. the endpoint it was created from or received on. It can not be reused and sent by another LINX endpoint.

If a process, which has allocated signal buffers, calls *fork(2)* , both instances will own copies of the signal buffers. As only one process may use a LINX endpoint and its resources, either the parent or the child MUST close the LINX endpoint using **linx_close**(3) .

## See Also

**linx**(7) , **linx_close**(3) , **linx_free_buf**(3) , **linx_hunt**(3) , **linx_send**(3) , **fork**(2)

## Author

Enea LINX team

# LINX_ATTACH(3) manual page

## Name

linx_attach() - Attach to and supervise a LINX endpoint
linx_detach() - Detach from an attached LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSATTREF linx_attach(LINX *** *linx* **, union LINX_SIGNAL ** *** *sig* **, LINX_SPID** *spid* **);**

**int linx_detach(LINX *** *linx* **, LINX_OSATTREF *** *attref* **);**

## Description

**linx_attach**() is used to supervise another LINX endpoint. When the other endpoint is closed or becomes unavailable, an attach signal is sent to the supervising LINX endpoint *linx* as a notification. If a signal *sig* is provided, this signal will be received when the attach is triggered. If *sig* is NULL, the default LINX attach signal with signal number LINX_OS_ATTACH_SIG is received instead. The linx_attach() call consumes the signal, taking over its ownership, and sets the *sig* pointer to LINX_NIL. The signal is consumed if an error occurs too. If *sig* is corrupt, **abort**(3) is called.

**linx_detach**() is used to detach from an attached LINX endpoint. It's an error to detach from a LINX endpoint more than once or to detach from a LINX endpoint after the attach signal has been received. It is not an error to detach if the attach signal is waiting in the receive queue of the LINX endpoint. To prevent multiple detaches, linx_detach() sets the attref pointer to LINX_ILLEGAL_ATTREF.

*linx* is the handle of the LINX endpoint.

*spid* is the identifier of the LINX endpoint to supervise.

*sig* is either NULL or a user defined LINX signal.

*attref* is the LINX_OSATTREF attach reference obtained from **linx_attach**().

## Return Value

**linx_attach**() returns an attach reference (LINX_OSATTREF) when successful. This attach reference can be used to cancel the attach (detach) after the attach is done. **linx_detach**() returns 0 on success.

On failure, **linx_attach**() returns LINX_ILLEGAL_ATTREF and **linx_detach**() returns -1, and *errno* is set appropriately.

## Errors

**EBADF**, **ENOTSOCK** The *linx* handle refers to an invalid socket descriptor.

# Bugs/Limitations

None.

# Example

In this example the server attaches to the client and tells the client to exit and waits for the attach signal.

```
Server:
#include <linx.h>
#define CLIENT_DONE_SIG 0x1234
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_hunt_sig[] = { 1, LINX_OS_HUNT_SIG };
  const LINX_SIGSELECT sel_att_sig [] = { 1, LINX_OS_ATTACH_SIG };
  /* Create a LINX endpoint with huntname "attacher" */
  linx = linx_open("attacher", NULL, 0);
  /* Hunt for the client */
  linx_hunt(linx, "client", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrive the clients spid */
  client = linx_sender(linx, &sig)
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
  /* Attach to the client */
  linx_attach(linx, client, NULL);
  /* Create "done" signal */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), CLIENT_DONE_SIG);
  /* Send "done" signal to client */
  linx_send(linx, &sig, client);
  /* Wait for the attach signal */
  linx_receive(linx, &sig, sel_att_sig);
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
Client:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_done_sig[] = { 1, CLIENT_DONE_SIG };

  /* Open a LINX endpoint with huntname "client" */
  linx = linx_open("client", NULL, 0);
  /* Wait for server to send "done" signal */
  linx_receive(linx, &sig, sel_done_sig);
  /* Close the LINX endpoint - this will trigger the attach */
  linx_close(linx);
  return 0;
}
```

# See Also

**linx**(7) , **linx_attach**(3) , **linx_close**(3) , **linx_detach**(3) , **linx_hunt**(3) , **linx_send**(3)

# Author

Enea LINX team

# LINX_CANCEL_NEW_LINK(3) manual page

## Name

linx_request_new_link() - request a new link signal.
linx_cancel_new_link() - cancel a new link request.

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_NLREF linx_request_new_link(LINX \***linx**, LINX_NLTOKEN** token**);**

**int linx_cancel_new_link(LINX \***linx**, LINX_NLREF \***nlref**);**

## Description

Request a new link signal from LINX when a new LINX link is available. The *token* is passed back and forth between the calling LINX endpoint and the LINX kernel module to keep track of which LINX link was sent in the last new link signal. The first time a LINX endpoint calls **linx_request_new_link**() the value of the *token* is ignored. The *token* is then returned in the new link signal from the LINX kernel module and should be used by the LINX endpoint in the next **linx_request_new_link**() call. If a LINX link unknown to the LINX endpoint is present when the **linx_request_new_link**() call is made a new link signal is immediately sent to the caller. Two or more consecutive calls to **linx_request_new_link**() are allowed from the same LINX endpoint. The **linx_request_new_link**() call returns a reference handle which is used in the **linx_cancel_new_link**() call if a request is to be cancelled. After a cancellation no more new link signals are received if the socket has no other pending new link requests.

*linx* is the handle of the own LINX endpoint.

*token* is passed back and forth between the calling LINX endpoint and the LINX kernel module.

*nlref* is the reference used to cancel a pending new link request.

The syntax of the new link signal received when a new link is available is:

```
struct linx_new_link {
    LINX_SIGSELECT signo;
    LINX_NLTOKEN token;
    int name;
    int attr;
    char buf[1];
};
```

**signo** is LINX_OS_NEW_LINK_SIG, **token** is the token value to be used in the next request for a new link signal, **name** is the offset into **buf** where the name of the new link is stored, the name is null terminated, **attr** is the offset into **buf** where the attributes, if any, of the link are stored, the attribute string is null terminted, **buf** is a character buffer containg the name and the attributes, if any, of the link.

# Return Value

**linx_request_new_link**() returnes an new link reference (LINX_NLREF) when successful, otherwise on failure LINX_ILLEGAL_NLREF is returned. **linx_cancel_new_link**() returns 0 on success and -1 on failure. In case of failure *errno* will be set.

# Errors

**EBADF** The *linx* handle refers to an invalid socket descriptor.

**ENOMEM** Insufficient memory is available.

**EINVAL** Returned by **linx_cancel_new_link**() when trying to cancel an non-existent new link request.

# Bugs/Limitations

None.

# See Also

**linx**(7)

# Author

Enea LINX team

# LINX_CANCEL_TMO(3) manual page

## Name

linx_request_tmo() - request timeout with specified signal
linx_cancel_tmo() - cancel the specified timeout.
linx_modify_tmo() - modify the specified timeout

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSTMOREF linx_request_tmo(LINX \***linx**, LINX_OSTIME** tmo**, union LINX_SIGNAL \*\***sig**);**

**int linx_cancel_tmo(LINX \***linx**, LINX_OSTMOREF \***tmoref**);**

**int linx_modify_tmo(LINX \***linx**, LINX_OSTMOREF \***tmoref**, LINX_OSTIME** tmo**);**

## Description

**linx_request_tmo**() is used to request a signal, *sig*, to be sent to the requesting LINX endpoint *linx* when a timeout has expired. Information necessary for cancelling the time-out is returned. The actual time-out time, *tmo*, is rounded upward to the next larger tick because a time-out can only trigger when a system clock tick is received. Hence, the routine guarantees at least the number of milliseconds requested, but may add a few more depending on the tick resolution. When several time-out signals are generated at the same clock tick, the order is unspecified. If *sig* is NULL, the default LINX timeout signal with signal number LINX_OS_TMO_SIG is received instead. The linx_request_tmo () call consumes the signal, and sets the *sig* pointer to LINX_NIL. The signal is also consumed if an error occurs.

**linx_cancel_tmo**() is used to cancel a timeout. The time-out is identified by the *tmoref* parameter, which was set by **linx_request_tmo**()**.** It's an error to cancel a timeout more than once or to cancel a timeout after the timeout signal has been received. It is not an error to cancel the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint. To prevent multiple cancellations, **linx_cancel_tmo**() sets the tmoref pointer to LINX_ILLEGAL_TMOREF.

**linx_modify_tmo**() is used to modify a timeout. The time-out is identified by the tmoref parameter, which was set by **linx_request_tmo**()**.** It's an error to modify a timeout after the timeout signal has been received. It is not an error to modify the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint.

*linx* is the handle of the LINX endpoint.

*spid* is the timeout time in milliseconds.

*sig* is either NULL or a user defined LINX signal.

*tmoref* is the reference to the timeout obtained from **linx_request_tmo**().

# Return Value

**linx_request_tmo**() returns an timeout reference (LINX_OSTMOREF) when successful. This timeout reference can be used to cancel the timeout. **linx_cancel_tmo**() returns 0 on success.

On failure, **linx_request_tmo**() returns LINX_ILLEGAL_TMOREF while **linx_cancel_tmo**() and **linx_modify_tmo**() both returns -1. In all cases *errno* will be set appropriately.

# Errors

**EBADF**, **ENOTSOCK** The *linx* handle refers to an invalid socket descriptor.

**ENOMEM** Not enough memory.

**EINVAL** Invalid argument.

# Bugs/Limitations

None.

# Example

This example shows how to request and wait for a timeout.

```
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_tmo_sig [] = { 1, LINX_OS_TMO_SIG };
  /* Create a LINX endpoint */
  linx = linx_open("tmo-test", NULL, 0);
  /* Request a one second timeout. When the timeout expires the
     default signal will be returned. */
  linx_request_tmo(linx, 1000, NULL);
  /* Wait for the timeout signal. */
  linx_receive(linx, &sig, sel_tmo_sig);
  /* Free the timeout signal */
  linx_free_buf(linx, &sig);
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
```

# See Also

**linx**(7) , **linx_close**(3) , **linx_open**(3) **linx_receive**(3) ,

# Author

Enea LINX team

# LINXCFG(1) manual page

## Name

linxcfg - a configuration utility for creating LINX links to other nodes

## Synopsis

**linxcfg [-t <cm>] [-v] <command> [COMMAND PARAMETERS and OPTIONS]**

**linxcfg [-t <cm>] [-v] destroy <linkname> [<link 2>...]**

**linxcfg -h**

**linxcfg help <cm>**

**linxcfg -t <cm> [help]**

## Description

**linxcfg** is the LINX configuration utility. Depending on the *command* given, **linxcfg** can destroy or create links using the specified Connection Manager (CM).

## Options

**linxcfg** has common options to all commands

which should be placed **before** the actual *command*

-h, --help
      Display help and exit.
-t, --type <cm>
      Specifies *cm* as the Connection Manager to use. If this option is not provided, *eth* will be used as the default CM.

      You can use **-h** to list all available CMs.
-v, --verbose
      Verbose mode.
-c, --cf, --conn-feature <feature[=value]>
      Adds a connection feature. This option can be specified several times.
-l, --lf, --link-feature <feature[=value]>
      Adds a link feature. This option can be specified several times.

## Commands

create PARAMETERS
      Creates a connection link to another node. The *PARAMETERS* depend on the CM specified, see below for more details.

      A remote link cannot be used until also the remote node creates the other end of the link (a link back to the originating node). The state is "connecting" until both sides are ready and communication via

the LINX protocol can begin.

destroy <linkname> [<link2>...]

> Destroy the connection link with the name *linkname*.

help [<cm>]

> Provides help about the given *cm* , either specified in **-t** or after the **help** command.

> If no *cm* was specified, it will show a list of available CMs.

# Options to Different Connection Managers

## Ethernet CM

The **Ethernet CM** is specified with the *eth* keyword, which is actually the default if no other CM was specified.

The **create** command takes the following options:

create <macaddr> <interface> <linkname> [OPTIONAL PARAMETERS]

macaddr

> Specify the remote mac address of the node to connect to, e.g. 0a:1b:2c:3d:4d:5e.

interface

> Specify the Ethernet interface to use. Ex.: eth0.

linkname

> Name of the connection link. This is the link part of the path in a remote hunt name to LINX
> endpoints on the other node. It will also be the name seen in */proc/net/linx/cm/eth/<linkname>*

Optional parameters

--window_size=nnn

> This is the send/receive window_size, in number of packets on the link, and may need to be modified
> to adapt LINX to really slow or really fast Ethernet performance.
> Both the local and remote side negotiate the window size, being the smallest of the two sizes the
> selected one. The default 128 messages should be sufficient in most configurations. The window size
> shall always be of a power of 2. Size 0 means to use the default window size.

--defer_queue_size=nnn

> The defer queue size is in packages, with the size depending on the LINX link. The defer queue is
> used on the sender side when the send queue is full. Every message sent, when the send queue is full,
> is stored in the defer queue, until the send queue has room for more messages. The defer queue size is
> the maximum number of packages, with messages, that can be stored in the defer queue before the
> link is disconnected, due to lack or resources. The default value 2048 packages should be sufficient in
> most systems. Size 0 means to use the default defer queue size.

--conn_tmo=nnn

> The connect timeout specifies the time, in milliseconds, to wait until a connection attempt fails and a
> new attempt is made. Once connected, it is used as a supervision time-out, i.e. every conn_tmo/3
> milliseconds, LINX checks if any packets from the peer has been received. LINX disconnects the
> connection after four consecutive conn_tmo/3 timeout periods without any packets from the peer.
> Tmo 0 means to use the default timeout (1000ms).

--mtu=nnn

> The MTU (Maximum Transmission Unit) specifies the size in bytes of the largets packet that the
> Ethernet protocol can pass onwards (this excludes the size of the Ethernet header). If not set the MTU
> is fetched from the interface. Typcially a MTU of 1500 bytes is used for Ethernet.

--attributes=s

> The attribute option gives the possibility to assign an arbitrary string to the link when it is created.
> This string is included in the new_link signal, which is sent to all link supervisors. More information
> on link supervisors can be found in

--coreid=nn

The coreid is used in a multicore environment to configure the destination coreid.

**linx**(7) man page.

# Tcp Cm

The TCP CM is specified using the *tcp* keyword.

The **create** command takes the following parameters:

create <ip> <linkname> [OPTIONAL PARAMETERS]

ip
>    The IP address to connect to.

linkname
>    Name of the connection link. This is the LINX link part of the path in a remote hunt name to LINX endpoints on the other node. It will also be the name seen in */proc/net/linx/cm/tcp/<linkname>*

Optional parameters

--live_tmo=<size>
>    The live_tmo parameter is the time in milliseconds between every heartbeat that is used to detect if the connection has gone down. The default value is 1000 ms.

--use_nagle=<bool>
>    Set to 1 if nagle algorithm shall be used on the socket for the connection. Default is off.

--attributes=<s>
>    The attribute option gives the possibility to assign an arbitrary string to the link when it is created. This string is included in the new_link signal, which is sent to all link supervisors.

# Rio Cm

The RIO CM is specified using the *rio* keyword.

The **create** command takes the following parameters:

create <local_port> <port> <dev_id> <mbox> <if> <link> [OPTIONAL PARAMETERS]

local_port
>    The local RapidIO port to connect to.

port
>    The remote RapidIO port to connect to.

dev_id
>    The RapidIO device to connect to.

mbox
>    The RapidIO mailbox id to use.

if
>    The RapidIO interface to use.

link
>    Name of the connection link. This is the LINX link part of the path in a remote hunt name to LINX endpoints on the other node. It will also be the name seen in */proc/net/linx/rlnh*

Optional parameters

--tmo=<tmo>
>    The tmo parameter is the time in hundreds of milliseconds between every heartbeat that is used to detect if the connection has gone down. The default value is 500 ms.

--mtu=<mtu>

Specifies the MTU (Maximum Transmission Unit) in bytes of the largest that the RIO CM can pass onwards, including the size of the RapidIO header. If not explicitly given, the MTU is fetched from the RapidIO device.

# Files

None.

# Diagnostics

*linxcfg* will display more verbose information to standard out, if the -v option is specified.

# Known Bugs

None.

# Examples

```
linxcfg create 01:23:a4:4f:b3:ac eth0 link_A
linxcfg destroy link_A
linxcfg -t tcp create 192.168.1.1 link_A
linxcfg -t tcp destroy link_A
linxcfg -t rio create 0 0 0 0 rio0 riolink0 --mtu=128 --tmo=10
linxcfg -t rio destroy riolink0
```

# See Also

**linx**(7) , **linxstat**(1) , **linxdisc**(8)

# Author

Enea LINX team

# LINX_CLOSE(3) manual page

## Name

linx_close() - Close a LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_close(LINX \***_linx_**);**

## Description

Close a LINX endpoint that was created by a **linx_open**(3) call. All resources owned by this LINX endpoint are freed. When a LINX endpoint is closed, attach signals are sent to other LINX endpoints that have attached to it using **linx_attach**(3) . A LINX endpoint will be automatically closed if its owner process exits without calling **linx_close**(3) first. Note that the LINX socket is not closed until the last reference has been closed. See **linx_open**(3) for more information about forking processes that own LINX endpoints.

_linx_ is the handle of the LINX endpoint to close.

## Return Value

On success, zero is returned. On error, -1 is returned and _errno_ is set.

## Errors

**EBADF** The _linx_ handle refers to an invalid socket descriptor.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_open**(3) , **linx_alloc**(3) , **linx_attach**(3)

## Author

Enea LINX team

# LINX_DETACH(3) manual page

## Name

linx_attach() - Attach to and supervise a LINX endpoint
linx_detach() - Detach from an attached LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSATTREF linx_attach(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *spid***);**

**int linx_detach(LINX \****linx***, LINX_OSATTREF \****attref***);**

## Description

**linx_attach**() is used to supervise another LINX endpoint. When the other endpoint is closed or becomes unavailable, an attach signal is sent to the supervising LINX endpoint *linx* as a notification. If a signal *sig* is provided, this signal will be received when the attach is triggered. If *sig* is NULL, the default LINX attach signal with signal number LINX_OS_ATTACH_SIG is received instead. The linx_attach() call consumes the signal, taking over its ownership, and sets the *sig* pointer to LINX_NIL. The signal is consumed if an error occurs too. If *sig* is corrupt, **abort**(3) is called.

**linx_detach**() is used to detach from an attached LINX endpoint. It's an error to detach from a LINX endpoint more than once or to detach from a LINX endpoint after the attach signal has been received. It is not an error to detach if the attach signal is waiting in the receive queue of the LINX endpoint. To prevent multiple detaches, linx_detach() sets the attref pointer to LINX_ILLEGAL_ATTREF.

*linx* is the handle of the LINX endpoint.

*spid* is the identifier of the LINX endpoint to supervise.

*sig* is either NULL or a user defined LINX signal.

*attref* is the LINX_OSATTREF attach reference obtained from **linx_attach**().

## Return Value

**linx_attach**() returns an attach reference (LINX_OSATTREF) when successful. This attach reference can be used to cancel the attach (detach) after the attach is done. **linx_detach**() returns 0 on success.

On failure, **linx_attach**() returns LINX_ILLEGAL_ATTREF and **linx_detach**() returns -1, and *errno* is set appropriately.

## Errors

**EBADF**, **ENOTSOCK** The *linx* handle refers to an invalid socket descriptor.

# Bugs/Limitations

None.

# Example

In this example the server attaches to the client and tells the client to exit and waits for the attach signal.

```
Server:
#include <linx.h>
#define CLIENT_DONE_SIG 0x1234
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_hunt_sig[] = { 1, LINX_OS_HUNT_SIG };
  const LINX_SIGSELECT sel_att_sig [] = { 1, LINX_OS_ATTACH_SIG };
  /* Create a LINX endpoint with huntname "attacher" */
  linx = linx_open("attacher", NULL, 0);
  /* Hunt for the client */
  linx_hunt(linx, "client", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrive the clients spid */
  client = linx_sender(linx, &sig)
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
  /* Attach to the client */
  linx_attach(linx, client, NULL);
  /* Create "done" signal */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), CLIENT_DONE_SIG);
  /* Send "done" signal to client */
  linx_send(linx, &sig, client);
  /* Wait for the attach signal */
  linx_receive(linx, &sig, sel_att_sig);
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
Client:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_done_sig[] = { 1, CLIENT_DONE_SIG };

  /* Open a LINX endpoint with huntname "client" */
  linx = linx_open("client", NULL, 0);
  /* Wait for server to send "done" signal */
  linx_receive(linx, &sig, sel_done_sig);
  /* Close the LINX endpoint - this will trigger the attach */
  linx_close(linx);
  return 0;
}
```

# See Also

linx(7) , linx_attach(3) , linx_close(3) , linx_detach(3) , linx_hunt(3) , linx_send(3)

# Author

Enea LINX team

# LINXDISC(8) manual page

## Name

linxdisc - The LINX discovery daemon

## Synopsis

**linxdisc [-d] [-c** *config-file* **] [-r** *retries* **]**

## Description

**linxdisc** is the LINX discovery daemon, which automatically detects and creates connections (links) to other remote LINX nodes, which also is running **linxdisc**. The daemon periodically sends advertisements and waits for advertisements from remote nodes. The period is about 3 seconds. Connections are created to allowed remote nodes, when advertisements are received.

See the configuration file *linxdisc.conf*(5) for filter rules, limiting which interfaces to use and which remote nodes to allow. Communication links are only allowed to remote nodes which advertise the same LINX network cluster name as configured and with a node name fulfilling the configured allow/deny rules. The configured LINX cluster name and node name are used in advertisements sent from this node. The cluster name and node name must be unique.

The advertised node name is the suggested linkname to use for links to the advertised node. The linkname is used in huntname paths by applications in other nodes, when they hunt for LINX endpoints in this node.

linxdisc must be run as root and there can only be one instance of linxdisc running on each node. When linxdisc catches a SIGHUP, the configuration will be reread. Use SIGTERM to terminate linxdisc. When linxdisc terminates, all connections are automatically destroyed.

The linxdisc protocol is described in detail in the Linx Protocol document (see http://linx.sourceforge.net).

## Options

-d
> Run in debug mode. The program is interactive, not run as a daemon, and all progress and activity is logged to standard error.

-c config-file
> Use *config-file* instead of the default configuration file */etc/linxdisc.conf*.

## Files

*/etc/linxdisc.conf*

> The default configuration file. See *linxdisc.conf*(5) for further details.

## Diagnostics

*linxdisc* writes all activity and errors to the syslog, unless the -d option was given.

## Known Bugs and Limitations

linxdisc cannot apply parameters such as timeouts etc when creating links. Interfaces that are unavailable when starting linxdisc will not be used when they become available, linxdisc has to be restarted in that case.

## See Also

**linx**(7) , **linx_hunt**(3) , **linxdisc.conf**(5)

## Author

Enea LINX team

# LINXDISC.CONF(5) manual page

## Name

linxdisc.conf - custom settings for linxdisc

## Synopsis

**/etc/linxdisc.conf**

## Description

This file contains configuration variables for the linxdisc daemon. For non-configured variables, default values are used.

The linxdisc.conf must be located in the /etc directory, unless specified when linxdisc is started.

The variables are set using this semantics: VALUE="value1 value2 ..."

When the linxdisc daemon catches a SIGHUP, the configuration will be reread. IFACE, LINX_NETWORK_NAME and NODE_NAME will not be changed during runtime, while ALLOW and DENY will be applied immediately to close disallowed connections and when creating new connections.

## Variables

**IFACE**

>	specifies what network interface(s) to use. Only real devices are supported. Several interfaces can be specified. e.g. IFACE="eth0 eth1". Default is to use all available devices.

**LINX_NETWORK_NAME**

>	specifies the name of a LINX cluster. linxdisc only creates LINX communication links to nodes, advertising the same cluster name. All nodes must have the linxdisc daemon running. Default is the cluster name "undefined".

**NODE_NAME**

>	specifies a node name to advertise. It is used by other nodes to filter against the allow/deny filter chains. It is also used as the recommended local link name, when linxdisc in a remote node creates LINX communication links to this node, as a response to the advertisement. Default node name is the hostname. Example: NODE_NAME="server1".

**ALLOW**

>	chain. If specified, linxdisc will only accept connections to nodes with the specified advertised node names. Names are given in a string, separated by spaces. This variable can be set several times in the configuration file and the names will be concatenated to one ALLOW chain. Example: ALLOW="node1 node2".

**DENY**

>	chain. NOTE: This setting will not have any effect if ALLOW is set. If DENY is specified instead of ALLOW, linxdisc will allow LINX connections to any node, except those with advertised node names listed here. Names are given in a string, separated by spaces. This variable can be set several times in the configuration file and the names will be concatenated to one DENY chain.

**PARAM**

>	List of parameters to apply to all links created by linxdisc. For a description of parameters and allowed values, see linxcfg(1) . This variable can be set several times in this file and the specified parameters are concatenated to one PARAM chain.

# See Also

**linxdisc**(8)

# Author

Enea LINX team

# LINX_FREE_BUF(3) manual page

## Name

linx_free_buf() - Free a signal buffer owned by a LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_free_buf(LINX *** *linx* **, union LINX_SIGNAL ** *** *sig* **);**

## Description

Free a signal buffer, owned by the LINX endpoint, associated with the *linx* handle. Once the buffer is freed it may not be accessed again. To be owned by this LINX endpoint, the buffer has either been received by or allocated at this endpoint.

*linx* is the handle to the LINX endpoint, owning the *sig* signal buffer.

*sig* is the signal buffer. When the call returns, the pointer (*sig) is set to LINX_NIL to prevent further access.

If *sig* is corrupt, **abort**(3) is called.

## Return Value

On success, 0 is returned. The function does not return on error.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_open**(3) , **linx_receive**(3) ,

## Author

Enea LINX team

# LINX_FREE_NAME(3) manual page

## Name

linx_get_name() - Get the name of another LINX endpoint
linx_free_name() - Free the name returned by **linx_get_name**(3)

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_get_name(LINX \***linx**, LINX_SPID** spid**, char \*\***name**);**

**int linx_free_name(LINX \***linx**, char \*\***name**);**

## Description

Get the full hunt name of the LINX endpoint associated with the spid binary identifier (spid). The hunt name is a string which may include a path of link names if the LINX endpoint is located on a remote node, e.g. "LinkToA/LinkToB/LinkToC/EndpointName".

*linx* is the handle of the own LINX endpoint.

*spid* is the spid of the LINX endpoint for which the name is requested.

*name* is a pointer to a hunt name buffer on successful return, otherwise NULL. Unless name is NULL on return, the name buffer must be freed with **linx_free_name**(3) , when it is no longer needed.

## Return Value

On success, 0 is returned and *name* contains a pointer to a buffer containing the huntname. The user is responsible for freeing the buffer with **linx_free_name**(3) .
In case of failure, -1 will be returned by both calls *errno* will be set.

## Errors

**EBADF** The *linx* handle refers to an invalid socket descriptor.

**EINVAL** The name parameter is NULL.

**ENOMEM** Insufficient memory is available.

## Bugs/Limitations

None.

# See Also

linx(7) , linx_open(3) , malloc(3) , free(3)

# Author

Enea LINX team

# LINX_FREE_STAT(3) manual page

## Name

linx_get_stat() - Get the statistics of a LINX endpoint
linx_free_stat() - Free the statistics struct returned by **linx_get_stat**(3)

## Synopsis

**#include <linx_types.h>**
**#include <linx_ioctl.h>**
**#include <linx.h>**

**int linx_get_stat(LINX \***linx**, LINX_SPID** spid**, struct linx_info_stat \*\***stat**);**

**int linx_free_stat(LINX \***linx**, struct linx_info_stat \*\***stat**);**

## Description

Get the statistics of the LINX endpoint, associated with the local *spid* binary identifier if the LINX kernel module is compiled with SOCK_STAT=yes.

*linx* is the handle to the LINX endpoint, which is used to retrive the statistics.

*spid* is the identifier of the other LINX endpoint, for which the statistics is requested.

*stat* will on return point to an struct *linx_info_stat* buffer contining the statistics. Unless returned as NULL, the buffer must be freed by the user with **linx_free_stat**(3) , when no longer needed. For details on the *linx_info_stat* struct see **linx**(7) .

## Return Value

Returns 0 if successful, for **linx_get_stat**(3) the *stat* will contain a pointer to a buffer, allocated by LINX, containing the statistics. The user is responsible for freeing the buffer with **linx_free_stat**(3) .
In case of failure, -1 will be returned and *errno* will be set.

## Errors

**ENOSYS** if the LINX kernel module has not been compiled with SOCK_STAT=yes.

**ENOMEM** Insufficient memory is available.

**EBADF** The *linx* handle is associated with an invalid socket descriptor.

## Bugs/Limitations

None.

# See Also

linx(7) , linx_open(3)

# Author

Enea LINX team

# LINX_GET_DESCRIPTOR(3) manual page

## Name

linx_get_descriptor() - Get the socket descriptor associated with a LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_get_descriptor(LINX \****linx***);**

## Description

The **linx_get_descriptor**() call returns the socket descriptor associated with a LINX endpoint. This socket descriptor is usually not needed when using the LINX API.

The LINX socket descriptor must be used if the application wants to use the generic **select**(2) or **poll**(2) calls. The application can then use LINX API function calls with the corresponding LINX endpoint handle to take care of events reported for a LINX socket descriptor. Alternatively, the LINX socket interface can be used.

*linx* is the handle of the LINX endpoint for which the socket descriptor is retrieved.

## Return Value

Returns a socket descriptor.

## Bugs/Limitations

None.

## Notes

A LINX socket descriptor retrieved by **linx_get_descriptor**(3) must NOT be closed by calling **close**(2) . Close the LINX endpoint instead using **linx_close**(3) .

## See Also

**linx**(7) , **linx_close**(3) ,**linx_open**(3) , **select**(2) , **poll**(2)

## Author

Enea LINX team

# LINX_GET_NAME(3) manual page

## Name

linx_get_name() - Get the name of another LINX endpoint
linx_free_name() - Free the name returned by **linx_get_name**(3)

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_get_name(LINX \***linx**, LINX_SPID** spid**, char \*\***name**);**

**int linx_free_name(LINX \***linx**, char \*\***name**);**

## Description

Get the full hunt name of the LINX endpoint associated with the spid binary identifier (spid). The hunt name is a string which may include a path of link names if the LINX endpoint is located on a remote node, e.g. "LinkToA/LinkToB/LinkToC/EndpointName".

linx is the handle of the own LINX endpoint.

spid is the spid of the LINX endpoint for which the name is requested.

name is a pointer to a hunt name buffer on successful return, otherwise NULL. Unless name is NULL on return, the name buffer must be freed with **linx_free_name**(3) , when it is no longer needed.

## Return Value

On success, 0 is returned and name contains a pointer to a buffer containing the huntname. The user is responsible for freeing the buffer with **linx_free_name**(3) .
In case of failure, -1 will be returned by both calls errno will be set.

## Errors

**EBADF** The linx handle refers to an invalid socket descriptor.

**EINVAL** The name parameter is NULL.

**ENOMEM** Insufficient memory is available.

## Bugs/Limitations

None.

# See Also

**linx**(7) , **linx_open**(3) , **malloc**(3) , **free**(3)

# Author

Enea LINX team

# LINX_GET_SPID(3) manual page

## Name

linx_get_owner() - Get the pid of the process that owns a LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**pid_t linx_get_owner(LINX \*linx, LINX_SPID spid);**

## Description

Get the OS pid of the process that owns the LINX endpoint referred to by the LINX spid.

## Return Value

Returns the pid.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_open**(3)

## Author

Enea LINX team

# LINX_GET_SPID(3) manual page

## Name

linx_get_spid() - Get the binary identifier of LINX endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_SPID linx_get_spid(LINX *** *linx* **);**

## Description

Get the binary identifier (spid) of the LINX endpoint referred to by the linx LINX handle.

## Return Value

Returns the spid.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_open**(3)

## Author

Enea LINX team

# LINX_GET_STAT(3) manual page

## Name

linx_get_stat() - Get the statistics of a LINX endpoint
linx_free_stat() - Free the statistics struct returned by **linx_get_stat**(3)

## Synopsis

**#include <linx_types.h>**
**#include <linx_ioctl.h>**
**#include <linx.h>**

**int linx_get_stat(LINX \***linx**, LINX_SPID** spid**, struct linx_info_stat \*\***stat**);**

**int linx_free_stat(LINX \***linx**, struct linx_info_stat \*\***stat**);**

## Description

Get the statistics of the LINX endpoint, associated with the local *spid* binary identifier if the LINX kernel module is compiled with SOCK_STAT=yes.

*linx* is the handle to the LINX endpoint, which is used to retrive the statistics.

*spid* is the identifier of the other LINX endpoint, for which the statistics is requested.

*stat* will on return point to an struct *linx_info_stat* buffer contining the statistics. Unless returned as NULL, the buffer must be freed by the user with **linx_free_stat**(3) , when no longer needed. For details on the *linx_info_stat* struct see **linx**(7) .

## Return Value

Returns 0 if successful, for **linx_get_stat**(3) the *stat* will contain a pointer to a buffer, allocated by LINX, containing the statistics. The user is responsible for freeing the buffer with **linx_free_stat**(3) .
In case of failure, -1 will be returned and *errno* will be set.

## Errors

**ENOSYS** if the LINX kernel module has not been compiled with SOCK_STAT=yes.

**ENOMEM** Insufficient memory is available.

**EBADF** The *linx* handle is associated with an invalid socket descriptor.

## Bugs/Limitations

None.

# See Also

linx(7) , linx_open(3)

# Author

Enea LINX team

# LINX_GET_VERSION(3) manual page

## Name

linx_get_version() - Retrieve the LINX version

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_get_version(char \****buf***);**

## Description

**linx_get_version**() is used to retrieve the LINX version. If buf is not NULL, the version is also converted to a string, "major.minor.patch", and placed in the user-supplied buffer pointed to by buf (which contains at least 14 bytes).

## Return Value

On success, the version is returned. Bit 31 to 24 contains the major version, bit 23 to 8 contains the minor version and bit 7 to 0 contains the patch version.
On error, -1 is returned.

## Errors

No errors are defined.

## See Also

**linx**(7)

## Author

Enea LINX team

# GWCMD(1) manual page

## Name

linxgwcmd - Gateway Command tool example

## Synopsis

**linxgwcmd [OPTIONS]**

## Description

Gateway Command tool example, can be used to find and test gateway servers on a network.

## Options

-a <auth>
>        Use <auth> as authentication string, which should be in "user:passwd" form

-b <brc_addr>
>        Use <brc_addr> as broadcast string, which should be in "udp://*:<port>" form, where <port> should
>        be the port number. Default broadcast string is "udp://*:21768"

-c <name>
>        Use <name> as the client's name (default "gw_client")

-e[<n>][,<b>]
>        Echo test to an OSE Gateway use <n> as number of loops (default 10) and <b> for number of
>        bytes/chunk (default 4)

-h
>        Print help

-l[<t>][,<n>]
>        List Gateways use <t> as timeout value (default is 5 secs) and <n> for max items (default lists all)
>        -l25,1 => look for Gateways in 25 secs and list only the

>        first found

-p <proc>
>        Hunt for process <proc>

-s <url/name>
>        Connect to a OSE GateWay

>        server using either the servers url or its name

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

`linxgwcmd`

## See Also

**linxgws**(8)

## Author

Enea LINX team

# LINXGWS(1) manual page

## Name

linxgws - start the LINX Gateway Server

## Synopsis

**linxgws** [*CONFIG FILE*]
**linxgws** *OPTION*

## Description

*CONFIG* FILE is given on the command line, then linxgws tries to locate a configuration file in /etc/linxgws.conf

**--h|--help**
    display this help and exit
**--version**
    output version information and exit
**The LINX Gateway Server will use the LINX kernel module loaded when clients connect.**
**The logging is assured via syslog check /var/log/messages /var/log/syslog and /var/log/user depending of your Linux distribution.**

## Author

Enea LINX team

## Reporting Bugs

Report bugs to <linx@enea.com>.

## Examples

```
linxgws /path/to/my/config/file.conf
```

## Files

/etc/linxgws.conf

## See Also

**linxgws.conf**(5) , **linxgwcmd**(1)

# Author

Enea

LINX team

# Author

# LINXGWS.CONF(5) manual page

## Name

linxgws.conf - configuration file for LINX Gateway Server

## Synopsis

**/etc/linxgws.conf**

## Description

This file contains configuration variables for the LINX Gateway Server. For some variables default values can be used, but only for those specified later in the follwing sections.If the configuration file doesn't exist it must be created.

The linxgws.conf must be located in the /etc directory.

The variables are set using this semantics: VALUE="value..."

When the daemon receives a SIGHUP it shuts down all the auxiliary processes the configuration will be reread. All the variables in the configuration file can changed. All variables must be separated with a newline.

## Variables

**GATEWAY_NAME**
> specifies the name of the LINX Gateway Server . This name will be advertised on the broadcast UDP port to all clients. There is no default value and the provisioning of this variable is mandatory.

Example: GATEWAY_NAME=example_gateway

**INTERFACE_NAME**
> specifies what network interface to use. The interface must be configured with at least one IP address. It can use only one interface. And it will use the first IP address configured on this interface. There is no default value and the provisioning of this variable is mandatory.

Example: INTERFACE_NAME=eth0

**PUBLIC_PORT**
> specifies the TCP port number on which the LINX Gateway Server will accept connection. It mustn't be in the reserved port range. It the Gateway daemon will not be able to open the port will shutdown automatically. Make sure the port is not blocked by the firewall. If the variable is left empty a default value will be used. The default value is 16384

Example: PUBLIC_PORT=35700

**BROADCAST_PORT**

      specifies the UDP port number on which the LINX Gateway Server will announce the gateway name and the TCP port on that it accepts connection. If the variable is left empty a default value will be used. The default value is 21768

Example: BROADCAST_PORT=21444

# See Also

**linxgws**(8) , **linxgwcmd**(5)

# Author

Enea LINX team

# LINX.H(3) manual page

## Name

linx.h - The LINX API

## Synopsis

**#include <linx.h>**

## Description

This is the main header file for applications using the LINX library. This library provides functions to handle communication between LINX endpoints using a messaging API.
Inclusion of the <linx.h> header may also make all symbols from *<stdint.h>* , *<sys/types.h>* and *<linx_types.h>* visible.

The **linx_open()** function creates a LINX endpoint and returns a handle that should be used in all other LINX API calls.

**LINX *linx_open(const char ***name**, uint32_t *options**, void ***arg**);**

**int linx_close(LINX ***linx**);**

**int linx_get_descriptor(LINX ***linx**);**

**union LINX_SIGNAL *linx_alloc(LINX ***linx**, LINX_OSBUFSIZE *size**, LINX_SIGSELECT *sig_no**);**

**int linx_free_buf(LINX ***linx**, union LINX_SIGNAL ****sig**);**

**int linx_send(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_SPID *to**);**

**int linx_send_w_opt(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_SPID *from**, LINX_SPID *to**, int32_t ***taglist**);**

**int linx_send_w_s(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_SPID *from**, LINX_SPID *to**);**

**int linx_receive(LINX ***linx**, union LINX_SIGNAL ****sig**, const LINX_SIGSELECT ***sig_sel**);**

**int linx_receive_w_tmo(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_OSTIME *tmo**, const LINX_SIGSELECT ***sig_sel**);**

**int linx_receive_from(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_OSTIME *tmo**, const LINX_SIGSELECT ***sig_sel**, LINX_SPID *from**);**

**LINX_SPID linx_sender(LINX ***linx**, union LINX_SIGNAL ****sig**);**

**int linx_sigattr(const LINX ***linx**, const union LINX_SIGNAL ****sig**, uint32_t *attr**, void ****value**);**

**LINX_OSBUFSIZE linx_sigsize(LINX ***linx**, union LINX_SIGNAL ****sig**);**

**int linx_set_sigsize(LINX ***linx**, union LINX_SIGNAL ****sig**, LINX_OSBUFSIZE *sigsize**);**

**int linx_hunt(LINX \****linx***, const char \****name***, union LINX_SIGNAL \*\****hunt_sig***);**

**int linx_hunt_from(LINX \****linx***, const char \****name***, union LINX_SIGNAL \*\****hunt_sig***, LINX_SPID** *from***);**

**LINX_OSATTREF linx_attach(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *spid***);**

**int linx_detach(LINX \****linx***, LINX_OSATTREF \****attref***);**

**LINX_SPID linx_get_spid(LINX \****linx***);**

**int linx_get_name(LINX \****linx***, LINX_SPID** *spid***, char \*\****name***);**

**int linx_free_name(LINX \****linx***, char \*\****name***);**

**int linx_get_stat(LINX \****linx***, LINX_SPID** *spid***, struct linx_info_stat \*\****stat***);**

**int linx_free_stat(LINX \****linx***, struct linx_info_stat \*\****stat***);**

**LINX_OSTMOREF linx_request_tmo(LINX \****linx***, LINX_OSTIME** *tmo***, union LINX_SIGNAL \*\****sig***);**

**int linx_cancel_tmo(LINX \****linx***, LINX_OSTMOREF \****tmoref***);**

**int linx_modify_tmo(LINX \****linx***, LINX_OSTMOREF \****tmoref***, LINX_OSTIME** *tmo***);**

**LINX_NLREF linx_request_new_link(LINX \****linx***, LINX_NLTOKEN** *token***);**

**int linx_cancel_new_link(LINX \****linx***, LINX_NLREF \*** *nlref***);**

**int linx_get_version(char \****buf***);**

# Future Directions

The LINX API is still under development and may be subject to change in future releases.

# Known Bugs

None.

# Author

Enea LINX team

# See Also

**linx**(7) , **linx_open**(3) , **linx_types.h**(3) , **linx_close**(3) , **linx_get_descriptor**(3) , **linx_alloc**(3) , **linx_free_buf**(3) , **linx_send**(3) , **linx_send_w_opt**(3) , **linx_send_w_s**(3) , **linx_receive**(3) , **linx_receive_w_tmo**(3) , **linx_receive_from**(3) , **linx_sender**(3) , **linx_sigattr**(3) , **linx_sigsize**(3) , **linx_set_sigsize**(3) , **linx_hunt**(3) , **linx_hunt_from**(3) , **linx_attach**(3) , **linx_detach**(3) , **linx_get_spid**(3) , **linx_get_name**(3) , **linx_free_name**(3) , **linx_get_stat**(3) , **linx_free_stat**(3) , **linx_request_tmo**(3) , **linx_modify_tmo**(3) , **linx_cancel_tmo**(3) , **linx_request_new_link**(3) , **linx_cancel_new_link**(3) **linx_get_version**(3)

# LINX_HUNT(3) manual page

## Name

linx_hunt() - Search for a LINX endpoint matching a name

linx_hunt_from() - Search for a LINX endpoint matching a name and set another LINX endpoint to be the owner of the hunt.

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_hunt(LINX \****linx***, const char \****name***, union LINX_SIGNAL \*\****hunt_sig***);**

**int linx_hunt_from(LINX \****linx***, const char \****name***, union LINX_SIGNAL \*\****hunt_sig***, LINX_SPID** *from***);**

## Description

**linx_hunt**() and **linx_hunt_from**() are used to search for a LINX endpoint matching the provided hunt *name,* usually in order to be able to obtain its spid and communicate with it. The hunt name may optionally be prepended by a path of link names to reach a remote node on which the endpoint shall be found. When the endpoint has been found, the provided signal *hunt_sig* is sent from the found endpoint to the calling endpoint. If no hunt signal (NULL) is provided, the default LINX hunt signal of type LINX_OS_HUNT_SIG is used. The call consumes the signal, taking ownership of the signal, and sets the *hunt_sig* pointer to LINX_NIL to prevent further access. The signal is consumed also if an error occurs.

If there are multiple endpoints with the same name, it is not specified which of these endpoints is used to resolve the hunt call. When the hunt signal is received, the *linx_sender*(3) call can be used to get the spid of the found endpoint. Note that the returned spid is always local even though it may correspond to a LINX endpoint on a remote node. LINX will transparently route signals to their intended destination endpoints on remote nodes when needed.

A pending hunt will be automatically cancelled if the *linx* endpoint that owns it is closed, e.g. if the caller exits. In the case with **linx_hunt_from**(), the *from* parameter indicates a different owner of the pending hunt. If the *from* endpoint is closed while the hunt is pending, the hunt will be freed. Note that *from* does NOT receive the hunt signal.

*name* is a null-terminated string that shall exactly match a LINX endpoints name. If the endpoint is located on a remote node, the name of a link to that node (or names in a sequence of links) shall be included as the leading part of the *name* according to the following syntax:
"[<link_1>/][<link_2>/]...[<link_N>/]name"
Example: "LinkToA/LinkToB/EndpointName" where a path of two links are passed to reach the LINX endpoint.

*linx* is the handle to the LINX endpoint that will receive the hunt signal.

*hunt_sig* is the signal which will be sent to the caller when the hunt has been resolved. The sender of the signal will be the matching LINX endpoint found. If *hunt_sig* is NULL, a default LINX hunt signal with signal number LINX_OS_HUNT_SIG will be sent instead. If *hunt_sig* is corrupt, **abort**(3) is called.

*from* specifies a different owner of a possible pending hunt. If the endpoint, with spid *from*, is closed, the pending hunt will be cancelled.

# Return Value

Returns 0 if successful. -1 if an error occurred in which case errno is set.

# Errors

**ENOMEM** Insufficient memory is available.

**EBADF** The *linx* handle refers to an invalid socket descriptor.

**EFAULT** The .I name parameter is invalid.

**ECONNRESET** The *from* spid in the **linx_hunt_from**(3) refers to a LINX endpoint that has been closed (i.e. in state LINX_ZOMBIE).

**EPIPE** This error is reported at an attempt to do a **linx_hunt_from**(3) and the *from* spid is a LINX endpoint that is being closed as the call occurs.

# Bugs/Limitations

None.

# Example

In this example client1 hunts for client2 and client3 */

```
Client 1:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client2, client3;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_hunt_sig[] = { 1, LINX_OS_HUNT_SIG };
  /* Open a LINX endpoint with huntname "client1" */
  linx = linx_open("client1", NULL, 0);
  /* Hunt for client2 */
  linx_hunt(linx, "client2", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrieve client2's spid */
  client2 = linx_sender(linx, &sig);
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
   /* Hunt for client3 */
  linx_hunt(linx, "client3", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrieve client3's spid */
  client3 = linx_sender(linx, &sig);
  /* Free the hunt signal */
```

Description                                                                                            71

```
  linx_free_buf(linx, &sig);
  /* Do work, e.g. exchanges signals with client2 or client3 ... */
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
Client 2:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  /* Open a LINX endpoint with huntname "client2" */
  linx = linx_open("client2", NULL, 0);
  /* Do work.... */
  /* When done close the LINX endpoint */
  linx_close(linx);
  /* "client2" can no longer be hunted on */
  return 0;
}
Client 3:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  /* Open a LINX endpoint with huntname "client3" */
  linx = linx_open("client3", NULL, 0);
  /* Do work.... */
  /* When done close the LINX endpoint */
  linx_close(linx);
  /* "client3" can no longer be hunted on */
  return 0;
}
```

# See Also

**linx**(7) , **linx_attach**(3) , **linx_sender**(3)

# Author

Enea LINX team

Example                                                                72

# LINX_HUNT_FROM(3) manual page

## Name

linx_hunt() - Search for a LINX endpoint matching a name

linx_hunt_from() - Search for a LINX endpoint matching a name and set another LINX endpoint to be the owner of the hunt.

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_hunt(LINX \***linx**, const char \***name**, union LINX_SIGNAL \*\***hunt_sig**);**

**int linx_hunt_from(LINX \***linx**, const char \***name**, union LINX_SIGNAL \*\***hunt_sig**, LINX_SPID** from**);**

## Description

**linx_hunt**() and **linx_hunt_from**() are used to search for a LINX endpoint matching the provided hunt *name,* usually in order to be able to obtain its spid and communicate with it. The hunt name may optionally be prepended by a path of link names to reach a remote node on which the endpoint shall be found. When the endpoint has been found, the provided signal *hunt_sig* is sent from the found endpoint to the calling endpoint. If no hunt signal (NULL) is provided, the default LINX hunt signal of type LINX_OS_HUNT_SIG is used. The call consumes the signal, taking ownership of the signal, and sets the *hunt_sig* pointer to LINX_NIL to prevent further access. The signal is consumed also if an error occurs.

If there are multiple endpoints with the same name, it is not specified which of these endpoints is used to resolve the hunt call. When the hunt signal is received, the *linx_sender*(3) call can be used to get the spid of the found endpoint. Note that the returned spid is always local even though it may correspond to a LINX endpoint on a remote node. LINX will transparently route signals to their intended destination endpoints on remote nodes when needed.

A pending hunt will be automatically cancelled if the *linx* endpoint that owns it is closed, e.g. if the caller exits. In the case with **linx_hunt_from**(), the *from* parameter indicates a different owner of the pending hunt. If the *from* endpoint is closed while the hunt is pending, the hunt will be freed. Note that *from* does NOT receive the hunt signal.

*name* is a null-terminated string that shall exactly match a LINX endpoints name. If the endpoint is located on a remote node, the name of a link to that node (or names in a sequence of links) shall be included as the leading part of the *name* according to the following syntax:
"[<link_1>/][<link_2>/]...[<link_N>/]name"
Example: "LinkToA/LinkToB/EndpointName" where a path of two links are passed to reach the LINX endpoint.

*linx* is the handle to the LINX endpoint that will receive the hunt signal.

*hunt_sig* is the signal which will be sent to the caller when the hunt has been resolved. The sender of the signal will be the matching LINX endpoint found. If *hunt_sig* is NULL, a default LINX hunt signal with signal number LINX_OS_HUNT_SIG will be sent instead. If *hunt_sig* is corrupt, **abort**(3) is called.

*from* specifies a different owner of a possible pending hunt. If the endpoint, with spid *from*, is closed, the pending hunt will be cancelled.

# Return Value

Returns 0 if successful. -1 if an error occurred in which case errno is set.

# Errors

**ENOMEM** Insufficient memory is available.

**EBADF** The *linx* handle refers to an invalid socket descriptor.

**EFAULT** The .I name parameter is invalid.

**ECONNRESET** The *from* spid in the **linx_hunt_from**(3) refers to a LINX endpoint that has been closed (i.e. in state LINX_ZOMBIE).

**EPIPE** This error is reported at an attempt to do a **linx_hunt_from**(3) and the *from* spid is a LINX endpoint that is being closed as the call occurs.

# Bugs/Limitations

None.

# Example

In this example client1 hunts for client2 and client3 */

```
Client 1:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client2, client3;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_hunt_sig[] = { 1, LINX_OS_HUNT_SIG };
  /* Open a LINX endpoint with huntname "client1" */
  linx = linx_open("client1", NULL, 0);
  /* Hunt for client2 */
  linx_hunt(linx, "client2", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrieve client2's spid */
  client2 = linx_sender(linx, &sig);
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
   /* Hunt for client3 */
  linx_hunt(linx, "client3", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, sel_hunt_sig);
  /* Retrieve client3's spid */
  client3 = linx_sender(linx, &sig);
  /* Free the hunt signal */
```

Description                                                                                          74

```
  linx_free_buf(linx, &sig);
  /* Do work, e.g. exchanges signals with client2 or client3 ... */
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
Client 2:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  /* Open a LINX endpoint with huntname "client2" */
  linx = linx_open("client2", NULL, 0);
  /* Do work.... */
  /* When done close the LINX endpoint */
  linx_close(linx);
  /* "client2" can no longer be hunted on */
  return 0;
}
Client 3:
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  /* Open a LINX endpoint with huntname "client3" */
  linx = linx_open("client3", NULL, 0);
  /* Do work.... */
  /* When done close the LINX endpoint */
  linx_close(linx);
  /* "client3" can no longer be hunted on */
  return 0;
}
```

# See Also

**linx**(7) , **linx_attach**(3) , **linx_sender**(3)

# Author

Enea LINX team

Example                                                                                      75

# LINX_MODIFY_TMO(3) manual page

## Name

linx_request_tmo() - request timeout with specified signal
linx_cancel_tmo() - cancel the specified timeout.
linx_modify_tmo() - modify the specified timeout

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSTMOREF linx_request_tmo(LINX \****linx***, LINX_OSTIME** *tmo***, union LINX_SIGNAL \*\****sig***);**

**int linx_cancel_tmo(LINX \****linx***, LINX_OSTMOREF \****tmoref***);**

**int linx_modify_tmo(LINX \****linx***, LINX_OSTMOREF \****tmoref***, LINX_OSTIME** *tmo***);**

## Description

**linx_request_tmo**() is used to request a signal, *sig*, to be sent to the requesting LINX endpoint *linx* when a timeout has expired. Information necessary for cancelling the time-out is returned. The actual time-out time, *tmo*, is rounded upward to the next larger tick because a time-out can only trigger when a system clock tick is received. Hence, the routine guarantees at least the number of milliseconds requested, but may add a few more depending on the tick resolution. When several time-out signals are generated at the same clock tick, the order is unspecified. If *sig* is NULL, the default LINX timeout signal with signal number LINX_OS_TMO_SIG is received instead. The linx_request_tmo () call consumes the signal, and sets the *sig* pointer to LINX_NIL. The signal is also consumed if an error occurs.

**linx_cancel_tmo**() is used to cancel a timeout. The time-out is identified by the *tmoref* parameter, which was set by **linx_request_tmo**()**.** It's an error to cancel a timeout more than once or to cancel a timeout after the timeout signal has been received. It is not an error to cancel the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint. To prevent multiple cancellations, **linx_cancel_tmo**() sets the tmoref pointer to LINX_ILLEGAL_TMOREF.

**linx_modify_tmo**() is used to modify a timeout. The time-out is identified by the tmoref parameter, which was set by **linx_request_tmo**()**.** It's an error to modify a timeout after the timeout signal has been received. It is not an error to modify the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint.

*linx* is the handle of the LINX endpoint.

*spid* is the timeout time in milliseconds.

*sig* is either NULL or a user defined LINX signal.

*tmoref* is the reference to the timeout obtained from **linx_request_tmo**().

# Return Value

**linx_request_tmo**() returns an timeout reference (LINX_OSTMOREF) when successful. This timeout reference can be used to cancel the timeout. **linx_cancel_tmo**() returns 0 on success.

On failure, **linx_request_tmo**() returns LINX_ILLEGAL_TMOREF while **linx_cancel_tmo**() and **linx_modify_tmo**() both returns -1. In all cases *errno* will be set appropriately.

# Errors

**EBADF**, **ENOTSOCK** The *linx* handle refers to an invalid socket descriptor.

**ENOMEM** Not enough memory.

**EINVAL** Invalid argument.

# Bugs/Limitations

None.

# Example

This example shows how to request and wait for a timeout.

```
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_tmo_sig [] = { 1, LINX_OS_TMO_SIG };
  /* Create a LINX endpoint */
  linx = linx_open("tmo-test", NULL, 0);
  /* Request a one second timeout. When the timeout expires the
     default signal will be returned. */
  linx_request_tmo(linx, 1000, NULL);
  /* Wait for the timeout signal. */
  linx_receive(linx, &sig, sel_tmo_sig);
  /* Free the timeout signal */
  linx_free_buf(linx, &sig);
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
```

# See Also

**linx**(7) , **linx_close**(3) , **linx_open**(3) **linx_receive**(3) ,

# Author

Enea LINX team

---

# LINXNET(8) manual page

## Name

linxnet - LINX ethernet driver

## Synopsis

**linxnet [-d] [-c** *config-file* **]**

## Description

*linxnet* is a virtual ethernet driver that uses *linx(7)* as carrier for Ethernet packets. *linxnet* opens a tap interface in order to receive or transmit packets from/to other linux application. This packets are transmitted over the linx link to the remote node.

See the configuration file *linxnet.conf*(5) for setting up link name, name and address of the tap interface. The linx link through witch the packets are send to the remote node, must be created by other program.

linxnet must be run as root and there can only be one instance of linxnet running on each node. When linxnet catches a SIGHUP, the configuration will be reread. Use SIGTERM to terminate linxnet. The tap interface opened by linxnet is automatically removed when linxnet exits.

## Options

-d
> Run in debug mode. The program is interactive, not run as a daemon, and all progress and activity is logged to standard error.

-c config-file
> Use *config-file* instead of the default configuration file */etc/linxnet.conf*.

## Files

*/etc/linxnet.conf*

> The default configuration file. See *linxnet.conf*(5) for further details.

## Diagnostics

*linxnet* writes all activity and errors to the syslog, unless the -d option was given.

## Known Bugs and Limitations

*linxnet* keeps a lockfile in /var/run folder. At startup, linxnet looks for this file and if it exist refuse to start thus preventing multiple instances of linxnet. In some cases, for example after a system crash, this file might not be removed properly thus preventing linxnet to start. The workaround is to remove the file manually.

# See Also

**linx**(7) , **linxnet.conf**(5)

# Author

Enea LINX team

# LINXNET.CONF(5) manual page

## Name

linxnet.conf - custom settings for linxnet

## Synopsis

*/etc/linxnet.conf*

## Description

This file contains configuration variables for the *linxnet* daemon. For non-configured variables, default values are used.

*linxnet.conf* must be located in the /etc directory, unless another directory is specified with the **-c** flag on the command line when linxnet is started.

The variables are set using this semantics: VALUE="value1 value2 ..."

If *linxnet* catches SIGHUP, *linxnet.conf(5)* will be reread and the new configuration will be applied immediately.

## Variables

**ETH_ADDRESS**
> specifies the MAC address of the interface created by *linxnet* Default address is 00:00:00:00:00:01. Example: ETH_ADDRESS="00:11:22:33:44:55".

**LINK_NAME**
> specfies the name of the *linx(7)* link that will be used for communication with the remote linxnet manager. Default is "link". Example: LINK_NAME="link_to_rhost".

**INTERFACE_NAME**
> specifies the name of the interface that will be created by *linxnet* and will be used by the Linux IP-stack and other network clients. Default is "tap". Example: INTERFACE_NAME="tap0".

**MTU**
> specifies the MTU for tap interface. Default is 1500. Example: MTU="8000"

**IP**

> specifies IP address for the interface created by *linxnet.* Default is 192.168.1.1. Example: IP="192.168.1.2"

**MASK**
> specifies netmask for the interface created by *linxnet.* Default is 255.255.255.0. Example: MASK="255.255.0.0"

## See Also

**linx**(7) , **linxnet**(8)

## Author

Enea LINX team

# LINX_OPEN(3) manual page

## Name

linx_open - Create a LINX endpoint and give it a huntname.

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX \*linx_open(const char \****name***, uint32_t** *options***, void \****arg***);**

## Description

A process uses **linx_open**() to create a LINX endpoint and give it a huntname. The call returns a handle to the LINX endpoint to be used in subsequent LINX API calls. As part of the **linx_open**(3) call, any pending hunts for the *name* are resolved, whether it is a local hunt or a hunt from a remote node connected via LINX.

*name* is the local part of the huntname to find this LINX endpoint. It is a null-terminated string, which in Linux may be of any length, but there may be restrictions due to other implementations, when communicating with other systems. *name* does not need to be unique in the system, but if multiple LINX endpoints have the same name it is unspecified which one is used when a **linx_hunt**(3) call for the name is resolved. From another LINX endpoint, use **linx_hunt**(3) with the huntname, followed by **linx_receive**(3) and **linx_sender**(3) , to get the binary identifier (spid) to be used to communicate with this LINX endpoint. The local LINX binary identifier (spid) for this endpoint can be fetched with **linx_get_spid**(3) but is usually not needed by the owner process.

The *options* and *arg* parameters are reserved for future extensions and shall be set to 0 and NULL respectively.

A LINX endpoint is closed by calling **linx_close**(3) or when the calling process exits.
If the process owning the LINX endpoint calls **fork**(2) , either the parent or the child must close the LINX handle with **linx_close**(3) . This is because only one process may use a LINX endpoint and after a fork both instances will have copies of the LINX handle. In any case, the LINX endpoint will not be terminated unless closed by both parent and child.

It is possible to retrieve the LINX socket descriptor associated with a LINX endpoint using **linx_get_descriptor**(3) if needed.

## Return Value

On success, a LINX handle is returned. On error, NULL is returned and *errno* is set.

## Errors

**EINVAL** Invalid argument passed.

**ENOBUFS** or **ENOMEM** Insufficient memory is available.

**EPROTONOSUPPORT** or **EAFNOSUPPORT** if *linx(7)* is not supported by the system.

**ENFILE** Not enough kernel memory to allocate a new LINX endpoint structure.

**EMFILE** Process file table overflow.

**EACCES** Permission to create a LINX endpoint is denied.

# Bugs/Limitations

None.

# See Also

**linx**(7) , **linx_close**(3) , **linx_hunt**(3) , **linx_get_descriptor**(3) , **linx_receive**(3) , **linx_send**(3) , **linx_sender**(3) , **fork**(2)

# Author

Enea LINX team

# LINX_RECEIVE(3) manual page

## Name

linx_receive() - Receive a LINX signal
linx_receive_w_tmo() - Receive a LINX signal with timeout
linx_receive_from() - Receive a LINX signal, but only from a given endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_receive(LINX *</i>linx<i>, union LINX_SIGNAL ***sig**, const LINX_SIGSELECT ***sig_sel</i>)**;**

**int linx_receive_w_tmo(LINX *</i>linx<i>, union LINX_SIGNAL ***sig**, LINX_OSTIME *tmo*, const LINX_SIGSELECT ***sig_sel</i>)**;**

**int linx_receive_from(LINX *</i>linx<i>, union LINX_SIGNAL ***sig**, LINX_OSTIME *tmo*, const LINX_SIGSELECT ***sig_sel</i>, LINX_SPID *from</i>)**;**

## Description

Receives a LINX signal. The calls will block until a signal is received or, if applicable, a timeout has elapsed. The first signal in the receive queue that matches the *sig_sel* filter is returned to the user. This way the user may process signals in a different order than when they arrived at the LINX endpoint. Signals that have arrived but have not yet been received using these calls will stay in the receive queue.

**linx_receive**() will wait indefinitely for a signal that matches the *sig_sel* filter. *sig_sel* filters are described in **linx**(7) .

**linx_receive_w_tmo**() waits until the provided timeout, *tmo*, has elapsed or for a signal that matches the *sig_sel* filter. When the timeout has elapsed, zero is returned and *sig* is set to LINX_NIL instead of a pointer to a signal buffer.

**linx_receive_from**() works in the same way as **linx_receive_w_tmo**() except that it will only accept signals sent from the LINX endpoint, indicated by the spid *from*.

*linx* is the handle to the LINX endpoint, via which the signals are received.

*sig* is the received signal. The signal buffer is allocated by **linx_receive**(3) . Never use a preallocated signal buffer as it will be lost.

*sig_sel* is a select filter, defining which types of buffers to receive. It is a list of signal numbers with a leading count indicating the number of signal numbers in the list. If the first position is set to a negative count, all signal numbers except those listed will be received. Read more about select filters in **linx**(7) .

*tmo* is the maximum number of milliseconds to wait. The value 0 will result in a check for a signal matching *sig_sel* followed by an immediate return.

*from* is the spid of the other LINX endpoint to receive from. Before **linx_receive_from**() is used, it is important to attach to the other LINX endpoint with **linx_attach**(3) and also to include the attach signal

number in the *sig_sel* filter. If this is not done and the other LINX endpoint (with spid *from* ) is closed, this call will block the entire specified timeout.

# Return Value

Returns the size of the received signal in bytes if successful. If a signal was received, it can be found in *sig*, otherwise *sig* will be LINX_NIL. Returns 0 if **linx_receive_w_tmo**(3) was called and no message was received before the timeout.

If unsuccessful, -1 is returned and *errno* is set.

# Errors

**ENOMEM** if there is not enough memory.

**EBADF**, **EFAULT**, **ENOTCONN**, **ENOTSOCK** if the underlying *linx* structure contains a invalid socket descriptor.

**EINTR** the call was interrupted by a signal.

# Bugs/Limitations

None.

# Notes

In case the application needs to wait on multiple sockets of different kinds, the internal socket descriptor in a LINX endpoint can be fetched with **linx_get_descriptor***(3)* and used in a **select***(2)* or **poll***(2)* call to wake up if anything arrives at the LINX endpoint.

# Example

In this example the server sends four signals to the client and the client chooses to receive them in the order it wants to.

```
Signal file (example.sig):
#include <linx.h>
#define SIG_X 0x1
#define SIG_Y 0x2
#define SIG_Z 0x3
#define SIG_W 0x4
/* select filter is { number-of-signals, signal, signal, ... } */
static const LINX_SIGSELECT sigsel_any[]     = { 0 };
static const LINX_SIGSELECT sigsel_sig_x_z[] = { 2, SIG_X, SIG_Z };
static const LINX_SIGSELECT sigsel_sig_w[]   = { 1, SIG_W };
Server:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
```

```
  union LINX_SIGNAL *sig;
  /* Open a linx endpoint with huntname "server" */
  linx = linx_open("server", NULL, 0);
  /* Hunt for client */
  linx_hunt(linx, "client", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, LINX_OS_HUNT_SIG);
  /* Retrieve client's spid */
  client = linx_sender(linx, &sig);
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
  /*  Send four signals, they will be stored in the receive
   *  queue on the client in same order as sent but the
   *  client chooses in which order to retrieve them from
   *  the queue.
   */
  /* Send signal with signal number SIG_X */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_X);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Y */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Y);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Z */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Z);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_W */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_W);
  linx_send(linx, &sig, client);
  linx_free_buf(linx, &sig);
  /* Close the linx endpoint */
  linx_close (linx);
}
Client:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
  /* Open a linx endpoint with huntname "client" */
  linx = linx_open("client", NULL, 0);
  /* Check for signal SIG_W first */
  linx_receive(linx, &sig, sigsel_sig_w);
  /* Do work, sig->sig_no is SIG_W */
  /* Free signal when done */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_X */
  linx_free_buf(linx, &sig);
  /* Receive either SIG_X or SIG_Z from the receive queue. */
  linx_receive(linx, &sig, sigsel_sig_x_z);
  /* Do work, sig->sig_no is SIG_Z (SIG_X has been consumed) */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_Y */
  linx_free_buf(linx, &sig);
  linx_close (linx);
}
```

## See Also

Example                                                              86

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) , linx_get_descriptor(3) , poll(2) , select(2)

# Author

Enea LINX team

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) , linx_get_descriptor(3) , poll(2) , select(2)

See Also

# LINX_RECEIVE_FROM(3) manual page

## Name

linx_receive() - Receive a LINX signal
linx_receive_w_tmo() - Receive a LINX signal with timeout
linx_receive_from() - Receive a LINX signal, but only from a given endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_receive(LINX *** *linx* **, union LINX_SIGNAL ** *** *sig* **, const LINX_SIGSELECT *** *sig_sel* **);**

**int linx_receive_w_tmo(LINX *** *linx* **, union LINX_SIGNAL ** *** *sig* **, LINX_OSTIME** *tmo* **, const LINX_SIGSELECT *** *sig_sel* **);**

**int linx_receive_from(LINX *** *linx* **, union LINX_SIGNAL ** *** *sig* **, LINX_OSTIME** *tmo* **, const LINX_SIGSELECT *** *sig_sel* **, LINX_SPID** *from* **);**

## Description

Receives a LINX signal. The calls will block until a signal is received or, if applicable, a timeout has elapsed. The first signal in the receive queue that matches the *sig_sel* filter is returned to the user. This way the user may process signals in a different order than when they arrived at the LINX endpoint. Signals that have arrived but have not yet been received using these calls will stay in the receive queue.

**linx_receive**() will wait indefinitely for a signal that matches the *sig_sel* filter. *sig_sel* filters are described in **linx**(7) .

**linx_receive_w_tmo**() waits until the provided timeout, *tmo*, has elapsed or for a signal that matches the *sig_sel* filter. When the timeout has elapsed, zero is returned and *sig* is set to LINX_NIL instead of a pointer to a signal buffer.

**linx_receive_from**() works in the same way as **linx_receive_w_tmo**() except that it will only accept signals sent from the LINX endpoint, indicated by the spid *from*.

*linx* is the handle to the LINX endpoint, via which the signals are received.

*sig* is the received signal. The signal buffer is allocated by **linx_receive**(3) . Never use a preallocated signal buffer as it will be lost.

*sig_sel* is a select filter, defining which types of buffers to receive. It is a list of signal numbers with a leading count indicating the number of signal numbers in the list. If the first position is set to a negative count, all signal numbers except those listed will be received. Read more about select filters in **linx**(7) .

*tmo* is the maximum number of milliseconds to wait. The value 0 will result in a check for a signal matching *sig_sel* followed by an immediate return.

*from* is the spid of the other LINX endpoint to receive from. Before **linx_receive_from**() is used, it is important to attach to the other LINX endpoint with **linx_attach**(3) and also to include the attach signal

number in the *sig_sel* filter. If this is not done and the other LINX endpoint (with spid *from* ) is closed, this call will block the entire specified timeout.

# Return Value

Returns the size of the received signal in bytes if successful. If a signal was received, it can be found in *sig*, otherwise *sig* will be LINX_NIL. Returns 0 if **linx_receive_w_tmo**(3) was called and no message was received before the timeout.

If unsuccessful, -1 is returned and *errno* is set.

# Errors

**ENOMEM** if there is not enough memory.

**EBADF**, **EFAULT**, **ENOTCONN**, **ENOTSOCK** if the underlying *linx* structure contains a invalid socket descriptor.

**EINTR** the call was interrupted by a signal.

# Bugs/Limitations

None.

# Notes

In case the application needs to wait on multiple sockets of different kinds, the internal socket descriptor in a LINX endpoint can be fetched with **linx_get_descriptor**(3) and used in a **select**(2) or **poll**(2) call to wake up if anything arrives at the LINX endpoint.

# Example

In this example the server sends four signals to the client and the client chooses to receive them in the order it wants to.

```
Signal file (example.sig):
#include <linx.h>
#define SIG_X 0x1
#define SIG_Y 0x2
#define SIG_Z 0x3
#define SIG_W 0x4
/* select filter is { number-of-signals, signal, signal, ... } */
static const LINX_SIGSELECT sigsel_any[]     = { 0 };
static const LINX_SIGSELECT sigsel_sig_x_z[] = { 2, SIG_X, SIG_Z };
static const LINX_SIGSELECT sigsel_sig_w[]   = { 1, SIG_W };
Server:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
```

```
  union LINX_SIGNAL *sig;
  /* Open a linx endpoint with huntname "server" */
  linx = linx_open("server", NULL, 0);
  /* Hunt for client */
  linx_hunt(linx, "client", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, LINX_OS_HUNT_SIG);
  /* Retrieve client's spid */
  client = linx_sender(linx, &sig);
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
  /*  Send four signals, they will be stored in the receive
   *  queue on the client in same order as sent but the
   *  client chooses in which order to retrieve them from
   *  the queue.
   */
  /* Send signal with signal number SIG_X */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_X);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Y */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Y);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Z */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Z);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_W */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_W);
  linx_send(linx, &sig, client);
  linx_free_buf(linx, &sig);
  /* Close the linx endpoint */
  linx_close (linx);
}
Client:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
  /* Open a linx endpoint with huntname "client" */
  linx = linx_open("client", NULL, 0);
  /* Check for signal SIG_W first */
  linx_receive(linx, &sig, sigsel_sig_w);
  /* Do work, sig->sig_no is SIG_W */
  /* Free signal when done */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_X */
  linx_free_buf(linx, &sig);
  /* Receive either SIG_X or SIG_Z from the receive queue. */
  linx_receive(linx, &sig, sigsel_sig_x_z);
  /* Do work, sig->sig_no is SIG_Z (SIG_X has been consumed) */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_Y */
  linx_free_buf(linx, &sig);
  linx_close (linx);
}
```

# See Also

Example                                                                                          90

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) , linx_get_descriptor(3) , poll(2) , select(2)

# Author

Enea LINX team

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) , linx_get_descriptor(3) , poll(2) , select(2)

See Also

Enea LINX team

# LINX_RECEIVE_W_TMO(3) manual page

## Name

linx_receive() - Receive a LINX signal
linx_receive_w_tmo() - Receive a LINX signal with timeout
linx_receive_from() - Receive a LINX signal, but only from a given endpoint

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_receive(LINX \****linx***, union LINX_SIGNAL \*\****sig***, const LINX_SIGSELECT \****sig_sel***);**

**int linx_receive_w_tmo(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_OSTIME** *tmo***, const LINX_SIGSELECT \****sig_sel***);**

**int linx_receive_from(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_OSTIME** *tmo***, const LINX_SIGSELECT \****sig_sel***, LINX_SPID** *from***);**

## Description

Receives a LINX signal. The calls will block until a signal is received or, if applicable, a timeout has elapsed. The first signal in the receive queue that matches the *sig_sel* filter is returned to the user. This way the user may process signals in a different order than when they arrived at the LINX endpoint. Signals that have arrived but have not yet been received using these calls will stay in the receive queue.

**linx_receive**() will wait indefinitely for a signal that matches the *sig_sel* filter. *sig_sel* filters are described in **linx**(7) .

**linx_receive_w_tmo**() waits until the provided timeout, *tmo*, has elapsed or for a signal that matches the *sig_sel* filter. When the timeout has elapsed, zero is returned and *sig* is set to LINX_NIL instead of a pointer to a signal buffer.

**linx_receive_from**() works in the same way as **linx_receive_w_tmo**() except that it will only accept signals sent from the LINX endpoint, indicated by the spid *from*.

*linx* is the handle to the LINX endpoint, via which the signals are received.

*sig* is the received signal. The signal buffer is allocated by **linx_receive**(3) . Never use a preallocated signal buffer as it will be lost.

*sig_sel* is a select filter, defining which types of buffers to receive. It is a list of signal numbers with a leading count indicating the number of signal numbers in the list. If the first position is set to a negative count, all signal numbers except those listed will be received. Read more about select filters in **linx**(7) .

*tmo* is the maximum number of milliseconds to wait. The value 0 will result in a check for a signal matching *sig_sel* followed by an immediate return.

*from* is the spid of the other LINX endpoint to receive from. Before **linx_receive_from**() is used, it is important to attach to the other LINX endpoint with **linx_attach**(3) and also to include the attach signal

number in the *sig_sel* filter. If this is not done and the other LINX endpoint (with spid *from* ) is closed, this call will block the entire specified timeout.

# Return Value

Returns the size of the received signal in bytes if successful. If a signal was received, it can be found in *sig*, otherwise *sig* will be LINX_NIL. Returns 0 if **linx_receive_w_tmo**(3) was called and no message was received before the timeout.

If unsuccessful, -1 is returned and *errno* is set.

# Errors

**ENOMEM** if there is not enough memory.

**EBADF**, **EFAULT**, **ENOTCONN**, **ENOTSOCK** if the underlying *linx* structure contains a invalid socket descriptor.

**EINTR** the call was interrupted by a signal.

# Bugs/Limitations

None.

# Notes

In case the application needs to wait on multiple sockets of different kinds, the internal socket descriptor in a LINX endpoint can be fetched with **linx_get_descriptor**(3) and used in a **select**(2) or **poll**(2) call to wake up if anything arrives at the LINX endpoint.

# Example

In this example the server sends four signals to the client and the client chooses to receive them in the order it wants to.

```
Signal file (example.sig):
#include <linx.h>
#define SIG_X 0x1
#define SIG_Y 0x2
#define SIG_Z 0x3
#define SIG_W 0x4
/* select filter is { number-of-signals, signal, signal, ... } */
static const LINX_SIGSELECT sigsel_any[]     = { 0 };
static const LINX_SIGSELECT sigsel_sig_x_z[] = { 2, SIG_X, SIG_Z };
static const LINX_SIGSELECT sigsel_sig_w[]   = { 1, SIG_W };
Server:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
```

```
  union LINX_SIGNAL *sig;
  /* Open a linx endpoint with huntname "server" */
  linx = linx_open("server", NULL, 0);
  /* Hunt for client */
  linx_hunt(linx, "client", NULL);
  /* Receive hunt signal */
  linx_receive(linx, &sig, LINX_OS_HUNT_SIG);
  /* Retrieve client's spid */
  client = linx_sender(linx, &sig);
  /* Free the hunt signal */
  linx_free_buf(linx, &sig);
  /*  Send four signals, they will be stored in the receive
   *  queue on the client in same order as sent but the
   *  client chooses in which order to retrieve them from
   *  the queue.
   */
  /* Send signal with signal number SIG_X */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_X);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Y */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Y);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_Z */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_Z);
  linx_send(linx, &sig, client);
  /* Send signal with signal number SIG_W */
  sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), SIG_W);
  linx_send(linx, &sig, client);
  linx_free_buf(linx, &sig);
  /* Close the linx endpoint */
  linx_close (linx);
}
Client:
#include <linx.h>
#include "example.sig"
int
main (int argc, char *argv[])
{
  LINX *linx;
  LINX_SPID client;
  /* Open a linx endpoint with huntname "client" */
  linx = linx_open("client", NULL, 0);
  /* Check for signal SIG_W first */
  linx_receive(linx, &sig, sigsel_sig_w);
  /* Do work, sig->sig_no is SIG_W */
  /* Free signal when done */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_X */
  linx_free_buf(linx, &sig);
  /* Receive either SIG_X or SIG_Z from the receive queue. */
  linx_receive(linx, &sig, sigsel_sig_x_z);
  /* Do work, sig->sig_no is SIG_Z (SIG_X has been consumed) */
  linx_free_buf(linx, &sig);
  /* Receive the the first signal waiting in the receive queue */
  linx_receive(linx, &sig, sigsel_any);
  /* Do work, sig->sig_no is SIG_Y */
  linx_free_buf(linx, &sig);
  linx_close (linx);
}
```

# See Also

Example

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) ,
linx_get_descriptor(3) , poll(2) , select(2)

## Author

Enea LINX team

linx(7) , linx_hunt(3) , linx_send(3) , linx_sender(3) , linx_free_buf(3) , linx_alloc(3) ,
linx_get_descriptor(3) , poll(2) , select(2)

See Also

# LINX_REQUEST_NEW_LINK(3) manual page

## Name

linx_request_new_link() - request a new link signal.
linx_cancel_new_link() - cancel a new link request.

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_NLREF linx_request_new_link(LINX \***linx**, LINX_NLTOKEN** token**);**

**int linx_cancel_new_link(LINX \***linx**, LINX_NLREF \***nlref**);**

## Description

Request a new link signal from LINX when a new LINX link is available. The *token* is passed back and forth between the calling LINX endpoint and the LINX kernel module to keep track of which LINX link was sent in the last new link signal. The first time a LINX endpoint calls **linx_request_new_link**() the value of the *token* is ignored. The *token* is then returned in the new link signal from the LINX kernel module and should be used by the LINX endpoint in the next **linx_request_new_link**() call. If a LINX link unknown to the LINX endpoint is present when the **linx_request_new_link**() call is made a new link signal is immediately sent to the caller. Two or more consecutive calls to **linx_request_new_link**() are allowed from the same LINX endpoint. The **linx_request_new_link**() call returns a reference handle which is used in the **linx_cancel_new_link**() call if a request is to be cancelled. After a cancellation no more new link signals are received if the socket has no other pending new link requests.

*linx* is the handle of the own LINX endpoint.

*token* is passed back and forth between the calling LINX endpoint and the LINX kernel module.

*nlref* is the reference used to cancel a pending new link request.

The syntax of the new link signal received when a new link is available is:

```
struct linx_new_link {
    LINX_SIGSELECT signo;
    LINX_NLTOKEN token;
    int name;
    int attr;
    char buf[1];
};
```

**signo** is LINX_OS_NEW_LINK_SIG, **token** is the token value to be used in the next request for a new link signal, **name** is the offset into **buf** where the name of the new link is stored, the name is null terminated, **attr** is the offset into **buf** where the attributes, if any, of the link are stored, the attribute string is null terminted, **buf** is a character buffer containg the name and the attributes, if any, of the link.

# Return Value

**linx_request_new_link**() returnes an new link reference (LINX_NLREF) when successful, otherwise on failure LINX_ILLEGAL_NLREF is returned. **linx_cancel_new_link**() returns 0 on success and -1 on failure. In case of failure *errno* will be set.

# Errors

**EBADF** The *linx* handle refers to an invalid socket descriptor.

**ENOMEM** Insufficient memory is available.

**EINVAL** Returned by **linx_cancel_new_link**() when trying to cancel an non-existent new link request.

# Bugs/Limitations

None.

# See Also

**linx**(7)

# Author

Enea LINX team

# LINX_REQUEST_TMO(3) manual page

## Name

linx_request_tmo() - request timeout with specified signal
linx_cancel_tmo() - cancel the specified timeout.
linx_modify_tmo() - modify the specified timeout

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSTMOREF linx_request_tmo(LINX \***linx**, LINX_OSTIME** tmo**, union LINX_SIGNAL \*\***sig**);**

**int linx_cancel_tmo(LINX \***linx**, LINX_OSTMOREF \***tmoref**);**

**int linx_modify_tmo(LINX \***linx**, LINX_OSTMOREF \***tmoref**, LINX_OSTIME** tmo**);**

## Description

**linx_request_tmo**() is used to request a signal, *sig*, to be sent to the requesting LINX endpoint *linx* when a timeout has expired. Information necessary for cancelling the time-out is returned. The actual time-out time, *tmo*, is rounded upward to the next larger tick because a time-out can only trigger when a system clock tick is received. Hence, the routine guarantees at least the number of milliseconds requested, but may add a few more depending on the tick resolution. When several time-out signals are generated at the same clock tick, the order is unspecified. If *sig* is NULL, the default LINX timeout signal with signal number LINX_OS_TMO_SIG is received instead. The linx_request_tmo () call consumes the signal, and sets the *sig* pointer to LINX_NIL. The signal is also consumed if an error occurs.

**linx_cancel_tmo**() is used to cancel a timeout. The time-out is identified by the *tmoref* parameter, which was set by **linx_request_tmo**()**.** It's an error to cancel a timeout more than once or to cancel a timeout after the timeout signal has been received. It is not an error to cancel the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint. To prevent multiple cancellations, **linx_cancel_tmo**() sets the tmoref pointer to LINX_ILLEGAL_TMOREF.

**linx_modify_tmo**() is used to modify a timeout. The time-out is identified by the tmoref parameter, which was set by **linx_request_tmo**()**.** It's an error to modify a timeout after the timeout signal has been received. It is not an error to modify the timeout if the timeout signal is waiting in the receive queue of the LINX endpoint.

*linx* is the handle of the LINX endpoint.

*spid* is the timeout time in milliseconds.

*sig* is either NULL or a user defined LINX signal.

*tmoref* is the reference to the timeout obtained from **linx_request_tmo**().

# Return Value

**linx_request_tmo**() returns an timeout reference (LINX_OSTMOREF) when successful. This timeout reference can be used to cancel the timeout. **linx_cancel_tmo**() returns 0 on success.

On failure, **linx_request_tmo**() returns LINX_ILLEGAL_TMOREF while **linx_cancel_tmo**() and **linx_modify_tmo**() both returns -1. In all cases *errno* will be set appropriately.

# Errors

**EBADF**, **ENOTSOCK** The *linx* handle refers to an invalid socket descriptor.

**ENOMEM** Not enough memory.

**EINVAL** Invalid argument.

# Bugs/Limitations

None.

# Example

This example shows how to request and wait for a timeout.

```
#include <linx.h>
int
main (int argc, char *argv[])
{
  LINX *linx;
  union LINX_SIGNAL *sig;
  const LINX_SIGSELECT sel_tmo_sig [] = { 1, LINX_OS_TMO_SIG };
  /* Create a LINX endpoint */
  linx = linx_open("tmo-test", NULL, 0);
  /* Request a one second timeout. When the timeout expires the
     default signal will be returned. */
  linx_request_tmo(linx, 1000, NULL);
  /* Wait for the timeout signal. */
  linx_receive(linx, &sig, sel_tmo_sig);
  /* Free the timeout signal */
  linx_free_buf(linx, &sig);
  /* Close the LINX endpoint */
  linx_close(linx);
  return 0;
}
```

# See Also

**linx**(7) , **linx_close**(3) , **linx_open**(3) **linx_receive**(3) ,

# Author

Enea LINX team

---

# LINX_SEND(3) manual page

## Name

linx_send() - Send a LINX signal to a LINX endpoint (spid)

linx_send_w_s() - Send a LINX signal and specify a different LINX endpoint as sender

linx_send_w_opt() - Send a signal to a LINX endpoint (spid) with alternative

methods

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_send(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *to***);**

**int linx_send_w_s(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *from***, LINX_SPID** *to***);**

**int linx_send_w_opt(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *from***, LINX_SPID** *to***,**
**int32_t \****taglist***);**

## Description

Send a signal to the LINX endpoint identified by the *to* binary identifier (spid). **linx_send**(3) always consumes the signal buffer, even if the call fails, by freeing the memory and setting \*sig to LINX_NIL to prevent further access.

*linx* is the handle to the LINX endpoint, via which the signal is transmitted.

*sig* is the signal to send. The signal buffer must be owned by the transmitting LINX endpoint. It could have been allocated with **linx_alloc**(3) or received via the same LINX endpoint. If *sig* is corrupt, **abort**(3) is called.

*to* is the identifier (spid) of the recipient LINX endpoint, usually found as a result of a **linx_hunt**(3) . If the *to* spid refers to an endpoint that has been closed (i.e. in state LINX_ZOMBIE), -1 is returned and *errno* is set to ECONNRESET. If *to* does not apply to the specific method used in **linx_send_w_opt**(3) , the LINX_ILLEGAL_SPID define shall be used as its value.

*from* is the identifier (spid) of a different sender endpoint (instead of the transmitting endpoint). The recipient will see this spid as the sender of the signal. If the *from* spid refers to an endpoint that has been closed (i.e. in state LINX_ZOMBIE), -1 is returned and *errno* is set to ECONNRESET. If *from* does not apply to the specific method used in **linx_send_w_opt**(3) , the LINX_ILLEGAL_SPID define shall be used as its value.

*taglist* is an array of tags and values used with the call **linx_send_w_opt**(3) . Each tag has an associated value that directly follows it. See the example for details. If the *taglist* contains an invalid parameter, -1 is returned and *errno* is set to EINVAL. The *taglist* parameter list can contain the folllowing values.

**LINX_SIG_OPT_OOB**

The **LINX_SIG_OPT_OOB** tag is used to enable the OOB send method. The value of the **LINX_SIG_OPT_OOB** tag shall be 1.

An OOB signal is a signal that is transmitted, transported, forwarded and received ahead of in band signals. When the OOB signal is placed in the signal queue of the receiving process, it is placed ahead of in band signals but after OOB signals present in the signal queue.

If the OOB signal is sent inter node and the operating system on the receiving node is not supporting OOB, the signal is delivered as an in band signal.

When sending an OOB signal the *sig* parameter is the signal to be sent OOB. The signal buffer must be owned by the transmitting LINX endpoint. It could have been allocated with **linx_alloc**(3) or received.

When an OOB signal is sent using **linx_send_w_opt**(3) , a signal attribute is set on the signal, **LINX_SIG_ATTR_OOB**. The **linx_sigattr**(3) is used to test if the attribute is set or not.

**LINX_SIG_OPT_END**

Marks the end of the taglist and has no value. Must be at the end of the *taglist*.

# Return Value

Returns 0 if successful. On failure, -1 is returned and *errno* will be set. In any case, *sig* is set to LINX_NIL.

# Errors

**EBADF**, **ENOTSOCK** The LINX endpoint is associated with an invalid socket descriptor.

**ENOBUFS**, **ENOMEM** Insufficient memory is available.

**ECONNRESET** The destination spid refers to a LINX endpoint that has been closed (i.e. in state LINX_ZOMBIE). This situation is best avoided by using *linx_attach*(3) to supervise the receiving LINX endpoint. Closed spid are not reused until forced by spid instance counter overflow.

**EPIPE** This error is reported at an attempt to send to the spid of a LINX endpoint that is being closed as the call occurs.

**EINVAL** This error is reported if an unrecognized taglist parameter was passed in the *taglist*.

# Bugs/Limitations

None.

# Example

```
union LINX_SIGNAL *sig;
 LINX *linx;
 LINX_SPID to_spid;
 LINX_SPID from_spid;
 int32_t taglist[3];
 ...
 taglist[0]=LINX_SIG_OPT_OOB;
 taglist[1]=1;
```

```
taglist[2]=LINX_SIG_OPT_END;
sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), 100);
(void)linx_send_w_opt(linx, &sig, to_spid, from_spid, taglist);
...
```

## Notes

It is illegal to allocate/receive a signal buffer on one LINX endpoint and then send it using **linx_send**(3) from another LINX endpoint. In that case a new signal buffer needs to be allocated and the buffer contents moved to the new signal buffers prior to sending it.

```
union LINX_SIGNAL *sig1, *sig2;
LINX_OSBUFSIZE size;
LINX_SIGSELECT sel[] = {0};
LINX *linx1, *linx2;
LINX_SPID to_spid;
...
sig1 = linx_receive(linx1, &sig1, sel);
size = linx_sigsize(linx1, sig1);
sig2 = linx_alloc(linx2, size, sel);
memcpy((void)sig2, (void)sig1, size);
(void)linx_free_buf(linx1, &sig1);
(void)linx_send(linx2, &sig2, to_spid);
...
```

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_hunt**(3) , **linx_receive**(3) , **linx_receive_w_tmo**(3) , **linx_receive_from**(3) , **linx_sigattr**(3)

## Author

Enea LINX team

Example                                                                                     102

# LINX_SENDER(3) manual page

## Name

linx_sender() - Retrieve the sender of a LINX signal

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_SPID linx_sender(LINX *_linx_, union LINX_SIGNAL **_sig_);**

## Description

Get the binary identifier (spid) of the LINX endpoint which sent the signal _sig_.

_linx_ is the handle of the LINX endpoint on which _sig_ was received.

If _sig_ is corrupt, **abort**(3) is called.

## Return Value

The identifier (spid) of the LINX endpoint, which sent the signal _sig_. If the signal was never sent, **linx_sender**() will return the spid of the own LINX endpoint.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_receive**(3) , **linx_receive_from**(3) , **linx_receive_w_tmo**(3) ,

## Author

Enea LINX team

**ENEA**

# LINX_SEND_W_OPT(3) manual page

## Name

linx_send() - Send a LINX signal to a LINX endpoint (spid)

linx_send_w_s() - Send a LINX signal and specify a different LINX endpoint as sender

linx_send_w_opt() - Send a signal to a LINX endpoint (spid) with alternative

methods

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_send(LINX \***linx**, union LINX_SIGNAL \*\***sig**, LINX_SPID** to**);**

**int linx_send_w_s(LINX \***linx**, union LINX_SIGNAL \*\***sig**, LINX_SPID** from**, LINX_SPID** to**);**

**int linx_send_w_opt(LINX \***linx**, union LINX_SIGNAL \*\***sig**, LINX_SPID** from**, LINX_SPID** to**,**
**int32_t \***taglist**);**

## Description

Send a signal to the LINX endpoint identified by the *to* binary identifier (spid). **linx_send**(3) always consumes the signal buffer, even if the call fails, by freeing the memory and setting *sig to LINX_NIL to prevent further access.

*linx* is the handle to the LINX endpoint, via which the signal is transmitted.

*sig* is the signal to send. The signal buffer must be owned by the transmitting LINX endpoint. It could have been allocated with **linx_alloc**(3) or received via the same LINX endpoint. If *sig* is corrupt, **abort**(3) is called.

*to* is the identifier (spid) of the recipient LINX endpoint, usually found as a result of a **linx_hunt**(3) . If the *to* spid refers to an endpoint that has been closed (i.e. in state LINX_ZOMBIE), -1 is returned and *errno* is set to ECONNRESET. If *to* does not apply to the specific method used in **linx_send_w_opt**(3) , the LINX_ILLEGAL_SPID define shall be used as its value.

*from* is the identifier (spid) of a different sender endpoint (instead of the transmitting endpoint). The recipient will see this spid as the sender of the signal. If the *from* spid refers to an endpoint that has been closed (i.e. in state LINX_ZOMBIE), -1 is returned and *errno* is set to ECONNRESET. If *from* does not apply to the specific method used in **linx_send_w_opt**(3) , the LINX_ILLEGAL_SPID define shall be used as its value.

*taglist* is an array of tags and values used with the call **linx_send_w_opt**(3) . Each tag has an associated value that directly follows it. See the example for details. If the *taglist* contains an invalid parameter, -1 is returned and *errno* is set to EINVAL. The *taglist* parameter list can contain the folllowing values.

**LINX_SIG_OPT_OOB**

The **LINX_SIG_OPT_OOB** tag is used to enable the OOB send method. The value of the **LINX_SIG_OPT_OOB** tag shall be 1.

An OOB signal is a signal that is transmitted, transported, forwarded and received ahead of in band signals. When the OOB signal is placed in the signal queue of the receiving process, it is placed ahead of in band signals but after OOB signals present in the signal queue.

If the OOB signal is sent inter node and the operating system on the receiving node is not supporting OOB, the signal is delivered as an in band signal.

When sending an OOB signal the *sig* parameter is the signal to be sent OOB. The signal buffer must be owned by the transmitting LINX endpoint. It could have been allocated with **linx_alloc**(3) or received.

When an OOB signal is sent using **linx_send_w_opt**(3) , a signal attribute is set on the signal, **LINX_SIG_ATTR_OOB**. The **linx_sigattr**(3) is used to test if the attribute is set or not.

**LINX_SIG_OPT_END**
Marks the end of the taglist and has no value. Must be at the end of the *taglist*.

# Return Value

Returns 0 if successful. On failure, -1 is returned and *errno* will be set. In any case, *sig* is set to LINX_NIL.

# Errors

**EBADF**, **ENOTSOCK** The LINX endpoint is associated with an invalid socket descriptor.

**ENOBUFS**, **ENOMEM** Insufficient memory is available.

**ECONNRESET** The destination spid refers to a LINX endpoint that has been closed (i.e. in state LINX_ZOMBIE). This situation is best avoided by using *linx_attach*(3) to supervise the receiving LINX endpoint. Closed spid are not reused until forced by spid instance counter overflow.

**EPIPE** This error is reported at an attempt to send to the spid of a LINX endpoint that is being closed as the call occurs.

**EINVAL** This error is reported if an unrecognized taglist parameter was passed in the *taglist*.

# Bugs/Limitations

None.

# Example

```
union LINX_SIGNAL *sig;
 LINX *linx;
 LINX_SPID to_spid;
 LINX_SPID from_spid;
 int32_t taglist[3];
 ...
 taglist[0]=LINX_SIG_OPT_OOB;
 taglist[1]=1;
```

```
taglist[2]=LINX_SIG_OPT_END;
sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), 100);
(void)linx_send_w_opt(linx, &sig, to_spid, from_spid, taglist);
...
```

## Notes

It is illegal to allocate/receive a signal buffer on one LINX endpoint and then send it using **linx_send**(3) from another LINX endpoint. In that case a new signal buffer needs to be allocated and the buffer contents moved to the new signal buffers prior to sending it.

```
union LINX_SIGNAL *sig1, *sig2;
LINX_OSBUFSIZE size;
LINX_SIGSELECT sel[] = {0};
LINX *linx1, *linx2;
LINX_SPID to_spid;
...
sig1 = linx_receive(linx1, &sig1, sel);
size = linx_sigsize(linx1, sig1);
sig2 = linx_alloc(linx2, size, sel);
memcpy((void)sig2, (void)sig1, size);
(void)linx_free_buf(linx1, &sig1);
(void)linx_send(linx2, &sig2, to_spid);
...
```

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_hunt**(3) , **linx_receive**(3) , **linx_receive_w_tmo**(3) , **linx_receive_from**(3) , **linx_sigattr**(3)

## Author

Enea LINX team

Example                                                                 106

# LINX_SEND_W_S(3) manual page

## Name

linx_send() - Send a LINX signal to a LINX endpoint (spid)

linx_send_w_s() - Send a LINX signal and specify a different LINX endpoint as sender

linx_send_w_opt() - Send a signal to a LINX endpoint (spid) with alternative

methods

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_send(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *to***);**

**int linx_send_w_s(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *from***, LINX_SPID** *to***);**

**int linx_send_w_opt(LINX \****linx***, union LINX_SIGNAL \*\****sig***, LINX_SPID** *from***, LINX_SPID** *to***,**
**int32_t \****taglist***);**

## Description

Send a signal to the LINX endpoint identified by the *to* binary identifier (spid). **linx_send**(3) always
consumes the signal buffer, even if the call fails, by freeing the memory and setting \*sig to LINX_NIL to
prevent further access.

*linx* is the handle to the LINX endpoint, via which the signal is transmitted.

*sig* is the signal to send. The signal buffer must be owned by the transmitting LINX endpoint. It could have
been allocated with **linx_alloc**(3) or received via the same LINX endpoint. If *sig* is corrupt, **abort**(3) is called.

*to* is the identifier (spid) of the recipient LINX endpoint, usually found as a result of a **linx_hunt**(3) . If the *to*
spid refers to an endpoint that has been closed (i.e. in state LINX_ZOMBIE), -1 is returned and *errno* is set to
ECONNRESET. If *to* does not apply to the specific method used in **linx_send_w_opt**(3) , the
LINX_ILLEGAL_SPID define shall be used as its value.

*from* is the identifier (spid) of a different sender endpoint (instead of the transmitting endpoint). The recipient
will see this spid as the sender of the signal. If the *from* spid refers to an endpoint that has been closed (i.e. in
state LINX_ZOMBIE), -1 is returned and *errno* is set to ECONNRESET. If *from* does not apply to the
specific method used in **linx_send_w_opt**(3) , the LINX_ILLEGAL_SPID define shall be used as its value.

*taglist* is an array of tags and values used with the call **linx_send_w_opt**(3) . Each tag has an associated value
that directly follows it. See the example for details. If the *taglist* contains an invalid parameter, -1 is returned
and *errno* is set to EINVAL. The *taglist* parameter list can contain the folllowing values.

**LINX_SIG_OPT_OOB**

The **LINX_SIG_OPT_OOB** tag is used to enable the OOB send method. The value of the **LINX_SIG_OPT_OOB** tag shall be 1.

An OOB signal is a signal that is transmitted, transported, forwarded and received ahead of in band signals. When the OOB signal is placed in the signal queue of the receiving process, it is placed ahead of in band signals but after OOB signals present in the signal queue.

If the OOB signal is sent inter node and the operating system on the receiving node is not supporting OOB, the signal is delivered as an in band signal.

When sending an OOB signal the *sig* parameter is the signal to be sent OOB. The signal buffer must be owned by the transmitting LINX endpoint. It could have been allocated with **linx_alloc**(3) or received.

When an OOB signal is sent using **linx_send_w_opt**(3) , a signal attribute is set on the signal, **LINX_SIG_ATTR_OOB**. The **linx_sigattr**(3) is used to test if the attribute is set or not.

**LINX_SIG_OPT_END**
Marks the end of the taglist and has no value. Must be at the end of the *taglist*.

# Return Value

Returns 0 if successful. On failure, -1 is returned and *errno* will be set. In any case, *sig* is set to LINX_NIL.

# Errors

**EBADF**, **ENOTSOCK** The LINX endpoint is associated with an invalid socket descriptor.

**ENOBUFS**, **ENOMEM** Insufficient memory is available.

**ECONNRESET** The destination spid refers to a LINX endpoint that has been closed (i.e. in state LINX_ZOMBIE). This situation is best avoided by using *linx_attach*(3) to supervise the receiving LINX endpoint. Closed spid are not reused until forced by spid instance counter overflow.

**EPIPE** This error is reported at an attempt to send to the spid of a LINX endpoint that is being closed as the call occurs.

**EINVAL** This error is reported if an unrecognized taglist parameter was passed in the *taglist*.

# Bugs/Limitations

None.

# Example

```
union LINX_SIGNAL *sig;
 LINX *linx;
 LINX_SPID to_spid;
 LINX_SPID from_spid;
 int32_t taglist[3];
 ...
 taglist[0]=LINX_SIG_OPT_OOB;
 taglist[1]=1;
```

```
taglist[2]=LINX_SIG_OPT_END;
sig = linx_alloc(linx, sizeof(LINX_SIGSELECT), 100);
(void)linx_send_w_opt(linx, &sig, to_spid, from_spid, taglist);
...
```

# Notes

It is illegal to allocate/receive a signal buffer on one LINX endpoint and then send it using **linx_send**(3) from another LINX endpoint. In that case a new signal buffer needs to be allocated and the buffer contents moved to the new signal buffers prior to sending it.

```
union LINX_SIGNAL *sig1, *sig2;
LINX_OSBUFSIZE size;
LINX_SIGSELECT sel[] = {0};
LINX *linx1, *linx2;
LINX_SPID to_spid;
...
sig1 = linx_receive(linx1, &sig1, sel);
size = linx_sigsize(linx1, sig1);
sig2 = linx_alloc(linx2, size, sel);
memcpy((void)sig2, (void)sig1, size);
(void)linx_free_buf(linx1, &sig1);
(void)linx_send(linx2, &sig2, to_spid);
...
```

# See Also

**linx**(7) , **linx_alloc**(3) , **linx_hunt**(3) , **linx_receive**(3) , **linx_receive_w_tmo**(3) , **linx_receive_from**(3) , **linx_sigattr**(3)

# Author

Enea LINX team

Example                                                                                                           109

# LINX_SET_SIGSIZE(3) manual page

## Name

linx_sigsize() - Get the size of a LINX signal buffer in bytes
linx_set_sigsize() - Set the size of a LINX signal buffer in bytes

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSBUFSIZE linx_sigsize(LINX \***linx**, union LINX_SIGNAL \*\***sig**);**

**int linx_set_sigsize(LINX \***linx**, union LINX_SIGNAL \*\***sig**, LINX_OSBUFSIZE** sigsize**);**

## Description

**linx_sigsize**() gets the size of the signal buffer *sig*. The signal buffer must be owned by the LINX endpoint, referred to by *linx* but can either be a received buffer or a buffer, allocated by the **linx_alloc**(3) call on this LINX endpoint. If *sig* is corrupt, **abort**(3) is called.

**linx_set_sigsize**() resizes the signal buffer *sig* to *sigsize* bytes and returns 0 on success. An error is returned if the size is outside permitted limits. If the new *size* of the signal buffer *sig* is smaller than the original size, the data stored in signal buffer *sig* is truncated.

*linx* is the handle to the LINX endpoint, which owns the *sig* signal buffer.

*sigsize* is the new size in bytes of the signal buffer.

## Return Value

*linx_sigsize* returns the size in bytes of the signal buffer. *linx_set_sigsize* returns zero on success. If an error occurred, -1 is returned and *errno* is set.

## Errors

**EMSGSIZE** *sigsize* is less than 4.

**ENOMEM** Insufficient memory is available.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_receive**(3)

# Author

Enea LINX team

# LINX_SIGATTR(3) manual page

## Name

linx_sigattr() - Get an attribute of a LINX signal

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**int linx_sigattr(const LINX *`linx`, const union LINX_SIGNAL **`sig`, uint32_t `attr`, void **`value`);**

## Description

The **linx_sigattr**(3) API function is used to determine what attributes a received signal sig has, due to a **linx_send_w_opt**(3) LINX API call. The *value* pointer will be set to an attribute specific value. The *linx* parameter is the LINX handle used when receiving the signal.

*linx* is the handle to the LINX endpoint, which owns the *sig* signal buffer.

*attr* is the attribute of the *sig* signal buffer.

*value* is the returned value that corresponds to *attr* which in turn is the name of the attribute set on the *sig* signal buffer. *attr* can have the following values.

**LINX_SIG_ATTR_OOB**
> A signal has the **LINX_SIG_OPT_OOB** attribute with value set to *1* if it was sent using the alternative **linx_send_w_opt**(3) with **LINX_SIG_OPT_OOB**. If the attribute is not set, the returned value is *~0*.

> The OOB attribute has no affect on further communication using the signal, it is kept to make it possible for the receiving LINX endpoint to know if a signal is OOB or not. If the signal is sent to another LINX endpoint and the OOB attribute needs to be preserved, the **linx_send_w_opt**(3) call needs to be used for that signaling. If the received OOB signal is sent with **linx_send**(3) or **linx_send_w_s**(3) the OOB attribute is lost and the signal is transmitted as a in band signal.

## Return Value

Returns 0 if successful. On failure, -1 is returned and *errno* will be set.

## Errors

**EBADF**, **ENOTSOCK** The LINX endpoint is associated with an invalid socket descriptor.

## Bugs/Limitations

None.

## See Also

linx(7) , linx_send_w_opt(3)

## Author

Enea LINX team

# LINX_SIGSIZE(3) manual page

## Name

linx_sigsize() - Get the size of a LINX signal buffer in bytes
linx_set_sigsize() - Set the size of a LINX signal buffer in bytes

## Synopsis

**#include <linx_types.h>**
**#include <linx.h>**

**LINX_OSBUFSIZE linx_sigsize(LINX \***_linx_**, union LINX_SIGNAL \*\***_sig_**);**

**int linx_set_sigsize(LINX \***_linx_**, union LINX_SIGNAL \*\***_sig_**, LINX_OSBUFSIZE** _sigsize_**);**

## Description

**linx_sigsize**() gets the size of the signal buffer _sig_. The signal buffer must be owned by the LINX endpoint, referred to by _linx_ but can either be a received buffer or a buffer, allocated by the **linx_alloc**(3) call on this LINX endpoint. If _sig_ is corrupt, **abort**(3) is called.

**linx_set_sigsize**() resizes the signal buffer _sig_ to _sigsize_ bytes and returns 0 on success. An error is returned if the size is outside permitted limits. If the new _size_ of the signal buffer _sig_ is smaller than the original size, the data stored in signal buffer _sig_ is truncated.

_linx_ is the handle to the LINX endpoint, which owns the _sig_ signal buffer.

_sigsize_ is the new size in bytes of the signal buffer.

## Return Value

_linx_sigsize_ returns the size in bytes of the signal buffer. _linx_set_sigsize_ returns zero on success. If an error occurred, -1 is returned and _errno_ is set.

## Errors

**EMSGSIZE** _sigsize_ is less than 4.

**ENOMEM** Insufficient memory is available.

## Bugs/Limitations

None.

## See Also

**linx**(7) , **linx_alloc**(3) , **linx_receive**(3)

# Author

Enea LINX team

# Author

# LINXSTAT(1) manual page

## Name

linxstat - display specific information about LINX

## Synopsis

**linxstat [-a] [-b] [-f] [-h] [-H] [-q] [-s] [-S] [-t** *type* **] [-T] [-l** *length* **] [-n** *hunt_name* **]**

## Description

**linxstat** is a utility to display the status of the LINX communication system, including information about all local and known remote LINX endpoints and LINX links to remote nodes. The names, types, status and local identifiers are displayed as well as queues, attaches and hunts. The information is fetched from the LINX kernel module. The following columns can be output, in addition to a leading summary (compare the example):

*PID* is the id of the process owning the LINX endpoint or LINX link. As each remote endpoint are represented by a local LINX endpoint with a local identifier (spid), they are all owned by one common local process.

*spid* is the local LINX binary identifier for the LINX endpoint or LINX link.

*name* is the huntname of the LINX endpoint or LINX link. If a name is too long for the column, a "+" sign is displayed and the tailing part of the name.

*type* is the type of entry, i.e. *local* for a local LINX endpoint, *link* for a LINX link to another node, *remote* for a local representation of a remote LINX endpoint, *zombie* for a LINX endpoint which has been closed, *illegal* if an endpoint structure is illegal. *unknown* if linxstat cannot reach the LINX endpoint.

*state* is the state of the endpoint or link. The main states are *running* not waiting for anything, *recv* while waiting to receive signal buffers, *poll* while waiting for events with the *poll*(2) or *select*(2) system call, *unknown* if linxstat cannot get the state.

*queue* is the receive queue. This is the number of signal buffers delivered to the LINX endpoint, but not yet received.

*attach(from/to)* are the pending attach information. Both the number of attaches to each LINX endpoint from other endpoints and the number of attaches from each endpoint to other endpoints are displayed.

*hunt* are pending hunts owned by each LINX endpoint. It also includes pending hunts, sent by others but with this endpoint as sender.

*tmo* are pending timeouts for each LINX endpoint.

## Options

-a
      Display more verbose pending attach information.
-b

Batch mode, only display the number of open LINX endpoints.

-f

Print receive filters.

-h

Display a short help message.

-H

Display more verbose pending hunt information.

-q

Display more verbose receive queue information.

-s

Display less information (simple). Only endpoint-specific information are displayed, no summary, no title and no counter columns.

-S

If the LINX kernel module has been compiled with per-socket statistics (SOCK_STAT) then linxstat can display statistics for each socket.

-t type

Display only sockets of the specified type.

-T

Print pending timeouts.

-l length

Specify length of name to be printed, default 18, minimum 5

-n hunt_name

Print info

only for hunt_name

# Files

None.

# Diagnostics

*linxstat* writes all information to standard output.

# Known Bugs

linxstat queries are not atomic, the state of the system may change while linxstat is running.

# Example

This is a displayed list from linxstat.

```
 local:5 remote:1 link:1
 pending attach:2
 pending hunt:0
 pending timeout:0
 queued signals:0
  PID spid     name               type   state  queue attach(f/t) hunt tmo
24903 0x10003 linx16             link   -      0     1/0         0    0
24903 0x10004 +16/apitest_server remote -      0     0/0         0    0
24909 0x10001 apitest_client     local  run    0     0/1         0    0
24909 0x10024 +est_attach_parent local  poll   0     1/1         0    0
```

```
24909 0x10005 not bound         local  run   0     0/0           0     0
24958 0x10023 +test_attach_child local  recv  0     0/0           0     0
24963 0x10025 linxstat          local  run   0     0/0           0     0
```

With -a, more information is displayed after each line with pending attaches. Extracts from a table, created by linxstat -a:

```
 local:4 remote:1 link:1
 pending attach:1
 pending hunt:0
 pending timeout:0
 queued signals:0
 PID   spid    name               type   state  queue attach(f/t) hunt tmo
 25053 0x10002 linx16             link   -      0     1/0           0     0
        Attach info (from):
           [attref:0x00000400   from:0x00010001 signo:0       size: 0]
 25053 0x10003 +16/apitest_server remote -      0     0/0           0     0
 25059 0x10001 apitest_client     local  run    0     0/1           0     0
        Attach info (to):
           [attref:0x00000400   to:0x00010001 signo:0       size:  0]
 25059 0x10004 not bound          local  run    0     0/0           0     0
 25059 0x10014 +inx_receive_w_tmo local  poll   0     0/0           0     0
 25083 0x10015 linxstat           local  run    0     0/0           0     0
```

# See Also

**linx**(7) , **linx_attach**(3) , **linx_hunt**(3) , **linx_open**(3) , **linxdisc**(8) , **linxdisc.conf**(5) , **linxcfg**(1)

# Author

Enea LINX team

Example                                                                118

# LINX_TYPES.H(3) manual page

## Name

linx_types.h - data types for LINX

## Synopsis

**#include <linx_types.h>**

## Description

In addition to the definitions below, the **linx_types.h** includes the system headers **<stdint.h>**, **<stdlib.h>** and **<linux/types.h>**.

This **linx_types.h** header may be included automatically by the **linx.h** header file.

The header file contains the following definitions of LINX types and constants:

**LINX**
> The LINX handle, returned from **linx_open**(3) . It is used to access the LINX endpoint with all other LINX API fuctions.

**LINX_FALSE**
> defined to 0.

**LINX_TRUE**
> defined to 1.

**LINX_ILLEGAL_ATTREF**
> defined to ((LINX_OSATTREF)0). It can be returned instead of a valid LINX_OSATTREF value, e.g. from **linx_attach**(3) .

**LINX_ILLEGAL_SPID**
> defined to ((LINX_SPID)0). It can be returned instead of a valid LINX_SPID value, e.g. from **linx_get_sped**(3) .

**LINX_NIL**
> defined to ((union LINX_SIGNAL *) 0). This null pointer value indicates that a signal pointer no longer points to any buffer, e.g. after freeing the signal buffer.

**LINX_OS_ATTACH_SIG**
> defined to ((LINX_SIGSELECT)252). This is the reserved signal number of the default attach signal used if the application does not provide a signal with the **linx_attach**(3) call.

**LINX_OS_HUNT_SIG**
> defined to ((LINX_SIGSELECT)251). This is the reserved signal number of the default hunt signal used if the application does not provide a signal with the **linx_hunt**(3) call.

**LINX_OSATTREF**
> A uint32_t used for attach references.

**LINX_OSBOOLEAN**
> An int used for boolean operations.

**LINX_OSBUFSIZE**
> An ssize_t used for signal buffer sizes.

**LINX_OSTIME**
> A uint32_t used for time in micro-seconds (us).

**LINX_SIGSELECT**
> A uint32_t used for signal numbers.

**LINX_SPID**

A uint32_t used for LINX endpoint binary identifiers (also called spid).

**LINX_SIG_OPT_END**

A define used as endmark in the *taglist* parameter list, passed in the **linx_send_w_opt**(3) call.

**LINX_SIG_OPT_OOB**

A define used for OOB signaling in the *taglist* parameter list, passed in the **linx_send_w_opt**(3) call.

**LINX_SIG_ATTR_OOB**

Is an attribute of a received signal that was sent OOB. Defined to 1.

The following LINX definition is not explicitly done in this header file, but is used in the LINX API.

**union LINX_SIGNAL**

is an opaque structure, containing any signal buffer of any size, with a 32-bit signal number of type **LINX_SIGSELECT** in the first four bytes. This signal number can be accessed as the *sig_no* struct member of the signal buffer. The application allocates various sized signal buffers with **linx_alloc**(3) , which returns a **union LINX_SIGNAL \***.

# See Also

**linx**(7) , **linx_alloc**(3) , **linx_attach**(3) , **linx_get_spid**(3) , **linx_hunt**(3) , **linx_open**(3)

# Author

Enea LINX team

# MKCMCLCON(1) manual page

## Name

mkcmclcon - create LINX Connection Manager Control Layer connections.

## Synopsis

**mkcmclcon [OPTIONS] <connection>**

## Description

Create a LINX Connection Manager Control Layer connection to a remote node.

## Options

-h, --help
> Display help and exit.

-c, --connection
> What connection to be used for this CMCL connection. The CMCL is commonly used on top of another LINX Connection Manager. But, this other Connection Manager can be very thin as the CMCL provides connection handling and supervision.

-s, --silent
> Silent mode, only error messages are printed to stdout.

-t, --con_tmo
> Connection timeout is used for supervision of the connection. The CMCL will send three alive packets during this timeout. If it has not received an alive packet from the peer upon sending the fourth alive packet, it will consider the connection as disconnected. Default value is 1000 (ms).

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
Create a CMCL connection between node A and node B. In this example, node
A and
node B have created a LINX ethernet connection to each other (ethcm_A and

ethcm_B).
```

```
On node A,
mkcmclcon -c ethcm/ethcm_B node_B
On node B,
mkcmclcon -c ethcm/ethcm_A node_A
```

# See Also

**linx**(7) , **linxcfg**(1) , **linxstat**(1) , **linxdisc**(8) , **mklink**(1) , **mkethcon**(1) , **rmcmclcon**(1)

# Author

Enea LINX team

# MKETHCON(1) manual page

## Name

mkethcon - create LINX Ethernet connections.

## Synopsis

**mkethcon [OPTIONS] <connection>**

## Description

Create a LINX Ethernet connection to a remote node.

## Options

-h, --help
> Display help and exit.

Mandatory options

-m, --mac=MAC address
> Specify the remote MAC address of the node to connect to, e.g. 0a:1b:2c:3d:4d:5e.

-i, --if=interface
> Specify the Ethernet interface to use, e.g. eth0.

Optional options

--window_size=nnn
> This is the send/receive window_size, in number of packets on the link, and may need to be modified to adapt LINX to really slow or really fast Ethernet performance.
> Both the local and remote side negotiate the window size, being the smallest of the two sizes the selected one. The default 128 messages should be sufficient in most configurations. The window size shall always be of a power of 2. Size 0 means to use the default window size.

--defer_queue_size=nnn
> The defer queue size is in packages, with the size depending on the LINX link. The defer queue is used on the sender side when the send queue is full. Every message sent, when the send queue is full, is stored in the defer queue, until the send queue has room for more messages. The defer queue size is the maximum number of packages, with messages, that can be stored in the defer queue before the link is disconnected, due to lack or resources. The default value 2048 packages should be sufficient in most systems. Size 0 means to use the default defer queue size.

--conn_tmo=nnn
> The connect timeout specifies the time, in milliseconds, to wait until a connection attempt fails and a new attempt is made. Once connected, it is used as a supervision time-out, i.e. every conn_tmo/3 milliseconds, LINX checks if any packets from the peer has been received. LINX disconnects the connection after four consecutive conn_tmo/3 timeout periods without any packets from the peer. Tmo 0 means to use the default timeout (1000ms).

--coreid=nn
> The coreid is used in a multicore environment to configure the destination coreid.

--mtu=nnn
> The MTU (Maximum Transmission Unit) specifies the size in bytes of the largets packet that the Ethernet protocol can pass onwards (this excludes the size of the Ethernet header). If not set the MTU is fetched from the interface. Typcially a MTU of 1500 bytes is used for Ethernet.

# Files

None.

# Diagnostics

None.

# Known Bugs

None.

# Examples

```
mkethcon --mac=01:23:a4:4f:b3:ac --if=eth0 ethcon_A
mkethcon -m 01:23:a4:4f:b3:ac -i eth0 ethcon_A
```

# See Also

**linx**(7) , **linxcfg**(1) , **linxstat**(1) , **linxdisc**(8) , **mklink**(1) , **rmethcon**(1)

# Author

Enea LINX team

# MKLINK(1) manual page

## Name

mklink - create LINX links.

## Synopsis

**mklink [OPTIONS] <link>**

## Description

Create a LINX link to a remote node.

## Options

-h, --help
> Display help and exit.

-a, --attributes=<attributes>
> Specify an optional cookie that is returned unmodified in the new_link signal. The new_link signal is sent to link supervisor(s) when the link is available.

-c, --connection=<connection>
> Specify the connection(s) that this link should use. This option can be repeated if a link should use more than one connection. Note that the connection(s) must be created before they can be assigned to a link. The connection name is made up of two parts, see example. The first part identifies the CM and the second part is the name that was used in the CM specific create command.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
mkethcon --mac=01:23:a4:4f:b3:ac --if=eth0 ethcon_A
mklink --connection=ethcm/ethcon_A link_A
mkethcon --mac=01:23:a4:4f:b3:ac --if=eth0 ethcon_A
mktcpcon tcpcon_A
mklink -c ethcm/ethcon_A -c tcpcm/tcpcon_A link_A
```

# See Also

linx(7) , linxcfg(1) , mkethcon(1) , mktcpcon(1) , rmlink(1)

# Author

Enea LINX team

# MKRIOCON(1) manual page

## Name

mkriocon - create LINX RapidIO connections.

## Synopsis

**mkriocon [OPTIONS] <connection>**

## Description

Create a LINX RapidIO connection to a remote node.

## Options

-h, --help
> Display help and exit.

Mandatory options

-p, --port=<rio port>
> Specify the remote port number of the node to connect to, e.g. 0.

-l, --local-port=<rio port>
> Specify the local port number to use, e.g. 0.

-I, --id=<rio device id>
> Specify the RapidIO device id to connect to, e.g. 1.

-i, --if=<rio device>
> Specify the RapidIO interface to use, e.g. rio0.

Optional options

-t, --tmo=nn
> The connection supervision timeout specifies the time (in hundreds of msec) to wait until a connection is considered broken. Default is 5 (yielding a timeout of 500 msec) and should be sufficient in most systems. Tmo 0 means to use the default timeout.

-M, --mtu=nnn
> The MTU (Maximum Transmission Unit) specifies the size in bytes of the largets packet that the RapidIO protocol can pass onwards (this includes the size of the RapidIO header). If not set the MTU is fetched from the interface. Typically an MTU of 256 bytes is used for RapidIO.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
mkriocon --port=0 --local-port=0 --mbox=1 --id=1 --if=rio0 riocon_A
mkriocon -p 0 -l 0 -m 1 -I 1 -i rio0 riocon_A
```

## See Also

**linx**(7) , **linxcfg**(1) , **linxstat**(1) , **mklink**(1) , **rmriocon**(1)

## Author

Enea LINX team

# MKSHMCON(1) manual page

## Name

mkshmcon - create LINX shared memory connections.

## Synopsis

**mkshmcon [OPTIONS] <connection>**

## Description

Create a LINX shared memory connection to a remote node.

## Options

-h, --help
> Display help and exit.

-m, --mtu
> The maximum transfer unit option defines the number of bytes that one slot in a mailbox can hold. This option is also used as maximum receive unit (i.e. symmetric Rx and Tx areas).

-s, --silent
> Silent mode, only error messages are printed to stdout.

-t, --con_tmo
> Connection time-out is used as a supervision time-out, i.e. every conn_tmo milliseconds, the shared memory CM checks if any packets from the peer has been received. After two consequtive conn_tmo:s without any packets from the peer, the connection is disconnected.

-x, --side
> Swap Tx and Rx area. One side's Tx area is the other side's Rx area and vice versa. When a connection is setup, the -x option must be set on one of the sides.

-n, --slots
> Number of slots per mailbox. This option must be the same on both sides of a connection.

-b, --mbox
> Mailbox identifier. Both sides of a connection must use the same identifier.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

# Examples

```
Create a shared memory connection between CPU A and CPU B. Each
mailbox consists of 16 slots and each slot can hold 120 bytes.
The connection time-out is set to 1 sec, i.e. if no packets has
been received from the peer for 2-3 sec, the connection is
disconnected:
On CPU A,
mkshmcon -t 1000 -b 1 -n 16 -m 120 cpu_b
On CPU B,
mkshmcon -t 1000 -b 1 -n 16 -m 120 -x cpu_a
```

# See Also

**linx**(7) , **linxcfg**(1) , **linxstat**(1) , **linxdisc**(8) , **mklink**(1) , **rmshmcon**(1)

# Author

Enea LINX team

# MKTCPCON(1) manual page

## Name

mktcpcon - create LINX TCP/IP connections.

## Synopsis

**mktcpcon [OPTIONS] <connection>**

## Description

Create a LINX TCP/IP connection to a remote node.

## Options

-h, --help
>    Display help and exit.

-i, --ipaddr=<IP address>
>    The IP address to connect to, e.g. 192.168.0.12.

-I, --ipv6addr=<IPv6 address>%<iface>
>    The IPv6 address to connect to, e.g. fe80::218:8bff:fe1e:2eef, followed by % as a delimiter and then
>    the interface the connection should be created on.

-t, --live_tmo=<size>
>    The live_tmo parameter is the time in milliseconds between every heartbeat that is used to detect if
>    the connection has gone down. The default value is 1000 ms.

-n, --use_nagle=<bool>
>    Set to 1 if nagle algorithm shall be used on the socket for the connection. Default is off.

-s, --silent
>    Silent mode.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
mktcpcon --ipaddr=192.168.0.12 tcpcon_A
mktcpcon --ipv6addr=fe80::218:8bff:fe1e:2eef%eth0 tcpcon_A
```

## See Also

**linx**(7) , **linxcfg**(1) , **mklink**(1) , **rmtcpcon**(1)

## Author

Enea LINX

team

# RBLOG(3) manual page

## Name

rblog - write to ram ringbuffer log

## Synopsis

**#include <rblog.h>**

**int rblog(const char \****fmt***, ...);**

## Description

Write a new message into the ring buffer. Older messages successively gets overwritten when the writing wraps around.

The function parameters *fmt* and ... follows the same format as the parameters to printk, i.e. a format string and optional data parameters.

## Notes

A maximum of 1024 bytes including null terminator will be written.

This function is reentrant and can be called in atomic context.

## Return Value

rblog() returns the number of characters written to the log after format conversions.

# RBLOG_CREATE(3) manual page

## Name

rblog_create - initialize ram ringbuffer log

## Synopsis

**#include <rblog.h>**

**int rblog_create(size_t** *size***);**

## Description

Allocate memory (at least *size* bytes; the value is rounded up to fit a whole number of memory pages), register file for reading the log in /proc, and register SysRq keys for printing the log to console. The name of the /proc file used and which keys are use for printing can be selected with #defines in rblog.h. *size* must be at least 1.

## Return Value

rblog_create() returns 0 if successful, or a negative error value otherwise.

## Errors

**EPERM** The buffer is already initialized, or *size* is less than 1.

**ENOMEM** Memory allocation for the buffer failed, or not enough memory to create /proc file entry.

rblog_create() can also return any error value returned by register_sysrq_key().

# RBLOG_DESTROY(3) manual page

## Name

rblog_destroy - remove ring buffer

## Synopsis

**#include <rblog.h>**

**void rblog_destroy(void);**

## Description

rblog_destroy() unregisters SysRq keys and the /proc file registered by rblog_create(), and frees the memory allocated for the log. After rblog_destroy() is called, rblog_create() can be called again to create a new log buffer.

# RBLOG-EXTRACT(1) manual page

## Name

rblog-extract - extract ram ringbuffer log from ELF core file

## Synopsis

**rblog-extract**

*core-file*

## Description

rblog-extract searches the memory segments of the given ELF core file for valid rblog structs, and (optionally) saves the log contents to file. The ELF core file should have been produced by Kdump after a kernel panic on a kernel where an rblog is being used (see the section BACKGROUND below).

## Notes

rblog-extract can read both 32- and 64-bit ELF files; however, the width of pointers and 'int's must be the same on the crashed machine and the machine running rblog-extract. rblog-extract will detect if this is not the case and exit with an error.

rblog-extract will not run (or compile) on non-Linux platforms; it uses Linux's definition of ELF files and also the Linux 'stat' system call.

## Background

If a kernel has crashed it can be very desirable to save the memory image of the crashed environment. This gives the possibility to debug it afterwards to (hopefully) find the cause of the problem. The kexec and kdump tools can be used to load and execute a new kernel from an existing kernel. This second kernel, called a "capture kernel" or "panic kernel", is executed after the first (running) kernel has crashed. In the environment of the second kernel, the memory of the first kernel can be accessed.

This is an example of manual usage. You only need to have the kexec-tools package installed:

1.
    Add "crashkernel=64M@16M" option to your bootloader configuration and reboot.
2.
    Load the panic kernel with, for example:

        kexec -p /boot/vmlinuz-2.6.22.17-0.1-default \

            --initrd=/boot/initrd-2.6.22.17-0.1-default \

        --append="root=/dev/sda5 1 irqpoll maxcpus=1"

The command line should include the root file system (root=), the irqpoll and reset_devices options and should end with the runlevel you would like to have in the kdump environment (probably 1). Use "man kexec" to see all options.

Now, if we get a crash, the "capture" (or "panic") kernel will be executed, and an ELF core file named "vmcore" that contains the memory segments of the crashed kernel will be saved in /var/log/dump/$date$time for later analysis.

# Author

Enea LINX team

# RMCMCLCON(1) manual page

## Name

rmcmclcon - remove LINX Connection Manager Control Layer connections.

## Synopsis

**rmcmclcon [OPTIONS] <connection>...**

## Description

Remove LINX Connection Manager Control Layer connections. Make sure that the connection is not assigned to any link before it is removed (i.e. remove the link first).

## Options

-h, --help
      Display help and exit.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmcmclcon cmclcon_A (or rmcmclcon cmclcm/cmclcon_A)
rmcmclcon cmclcon_A cmclcon_B
```

## See Also

**linx**(7) , **linxcfg**(1) , **mkcmclcon**(1) , **rmlink**(1)

## Author

Enea LINX

team

# RMETHCON(1) manual page

## Name

rmethcon - remove LINX Ethernet connections.

## Synopsis

**rmethcon [OPTIONS] <connection>...**

## Description

Remove the LINX Ethernet connections. Make sure that the connection is not assigned to any link before it is removed (i.e. remove the link first).

## Options

-h, --help
        Display help and exit.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmethcon ethcon_A (or rmethcon ethcm/ethcon_A)
rmethcon ethcon_A ethcon_B
```

## See Also

**linx(7)** , **linxcfg(1)** , **mkethcon(1)** , **rmlink(1)**

## Author

Enea LINX

team

# RMLINK(1) manual page

## Name

rmlink - remove LINX links.

## Synopsis

**rmlink [-ahs] link...**

## Description

Remove LINX links and optionally any connections that they uses.

## Options

-a, --all
> Remove all, i.e. both the link and its assigned connection(s) are removed. This is the preferred method. If this option isn't used, any connections assigned to the link must be removed with a separate command, e.g. rmethcon, rmtcpcon, etc.

-h, --help
> Display help and exit.

-s, --silent
> Silent mode.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmlink -a link_A
rmlink link_A link_B
```

## See Also

**linx**(7) , **linxcfg**(1) , **mklink**(1) , **rmlink**(1)

# Author

Enea LINX

team

# RMRIOCON(1) manual page

## Name

rmriocon - remove LINX RapidIO connections.

## Synopsis

**rmriocon [OPTIONS] <connection>...**

## Description

Remove the LINX RapidIO connections. Make sure that the connection is not assigned to any link before it is removed (i.e. remove the link first).

## Options

-h, --help
Display help and exit.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmriocon riocon_A (or rmriocon riocm/riocon_A)
rmriocon riocon_A riocon_B
```

## See Also

**linx(7)** , **linxcfg(1)** , **mkriocon(1)** , **rmlink(1)**

## Author

Enea LINX

team

# RMSHMCON(1) manual page

## Name

rmshmcon - remove LINX shared memory connections.

## Synopsis

**rmshmcon [OPTIONS] <connection>...**

## Description

Remove LINX shared memory connections. Make sure that the connection is not assigned to any link before it is removed (i.e. remove the link first).

## Options

-h, --help
    Display help and exit.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmshmcon shmcon_A (or rmshmcon shmcm/shmcon_A)
rmshmcon shmcon_A shmcon_B
```

## See Also

**linx**(7) , **linxcfg**(1) , **mkshmcon**(1) , **rmlink**(1)

## Author

Enea LINX

team

# RMTCPCON(1) manual page

## Name

rmtcpcon - remove LINX TCP/IP connections.

## Synopsis

**rmtcpcon [OPTIONS] <connection>...**

## Description

Remove LINX TCP/IP connections. Make sure that the connection is not assigned to any link before it is removed (i.e. remove the link first).

## Options

-h, --help
        Display help and exit.

## Files

None.

## Diagnostics

None.

## Known Bugs

None.

## Examples

```
rmethcon tcpcon_A
rmethcon tcpcon_A tcpcon_B
```

## See Also

**linx**(7) , **linxcfg**(1) , **mktcpcon**(1) , **rmlink**(1)

## Author

Enea LINX

team

# Copyright